Data Processing at the Speed of 100 Gbps using Apache Crail

Patrick Stuedi
IBM Research

Apache Crail (crail.apache.org)

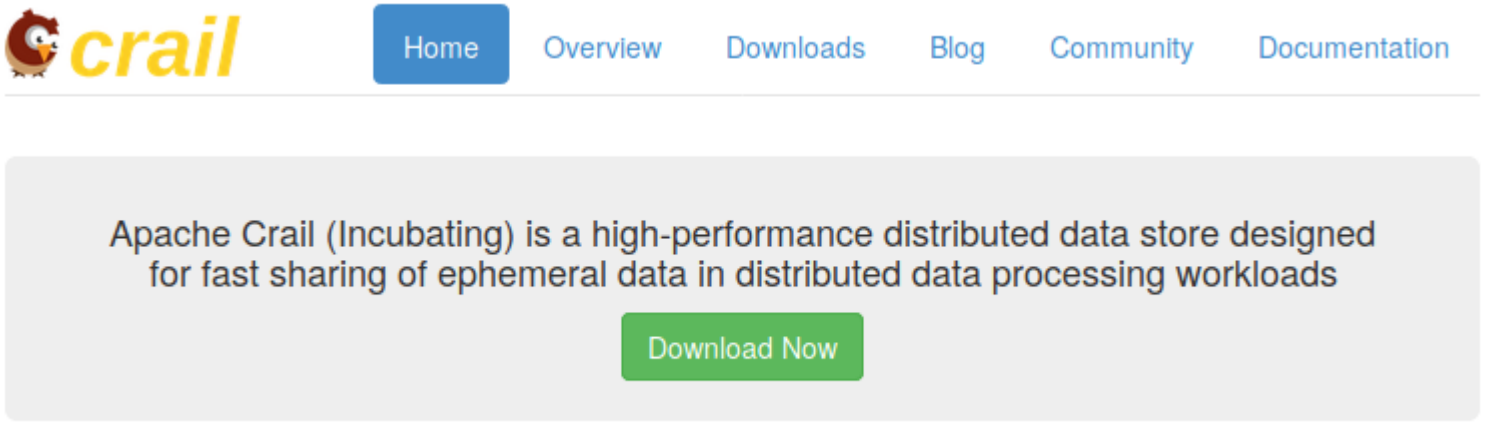

Apache Crail (crail.apache.org)

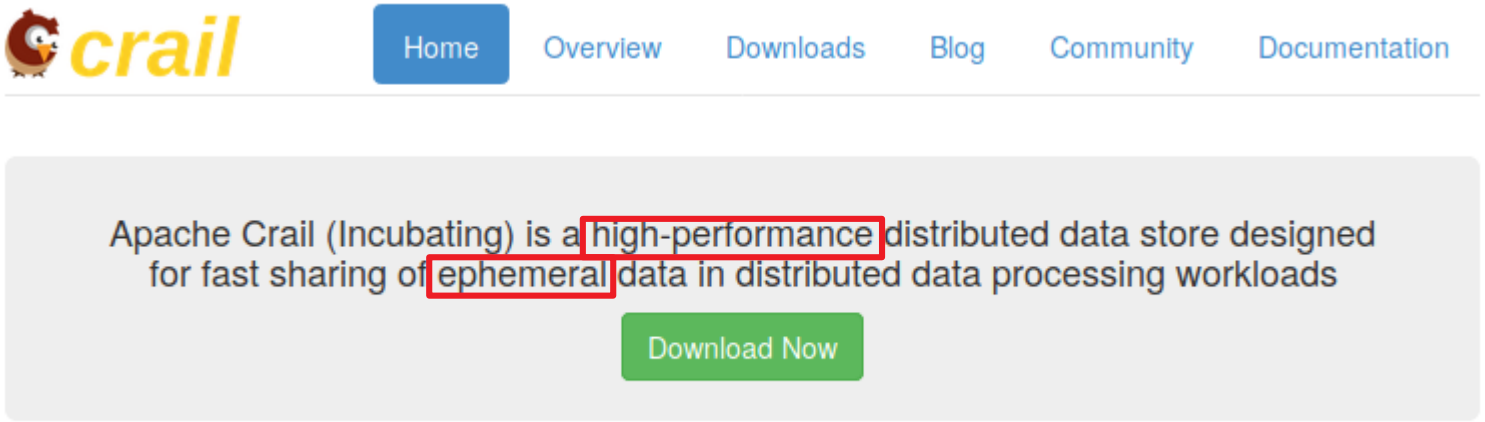

Ephemeral Data

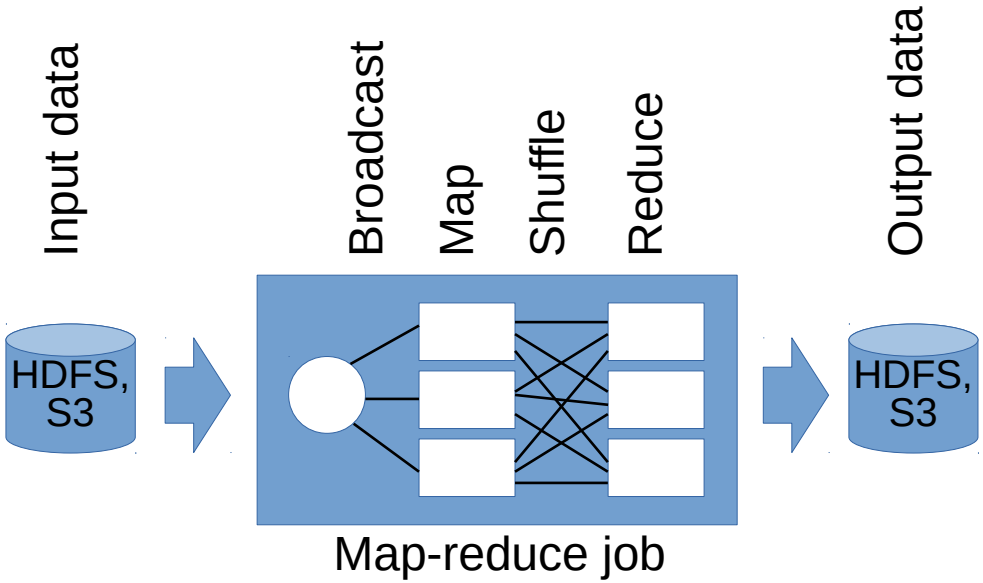

Ephemeral Data

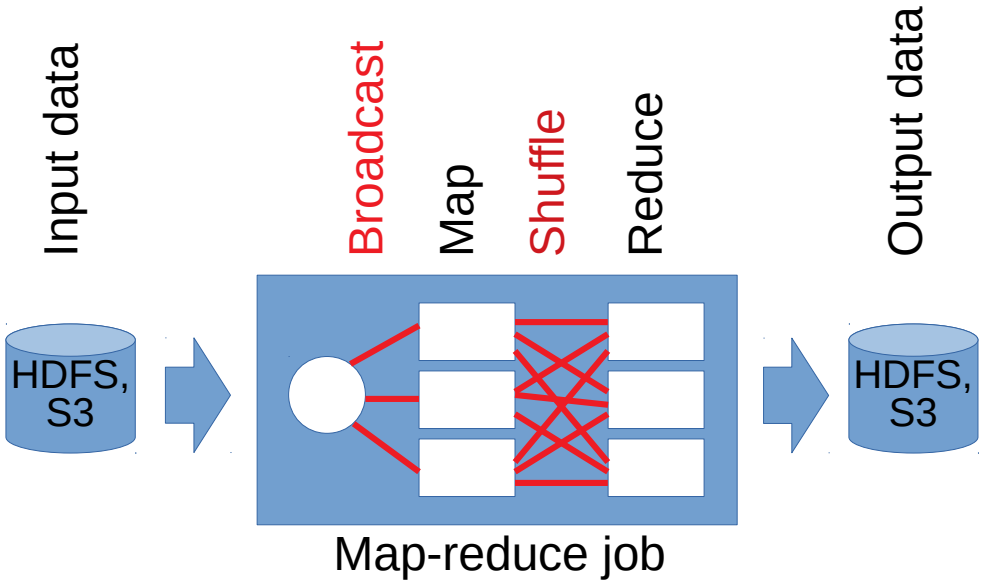

Ephemeral Data

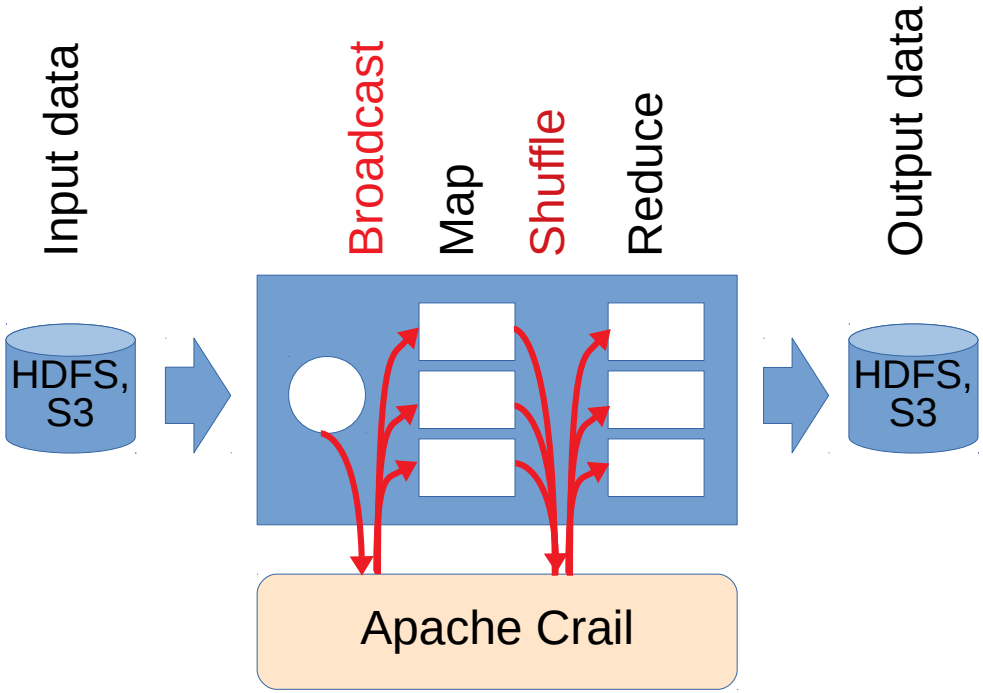

Ephemeral Data

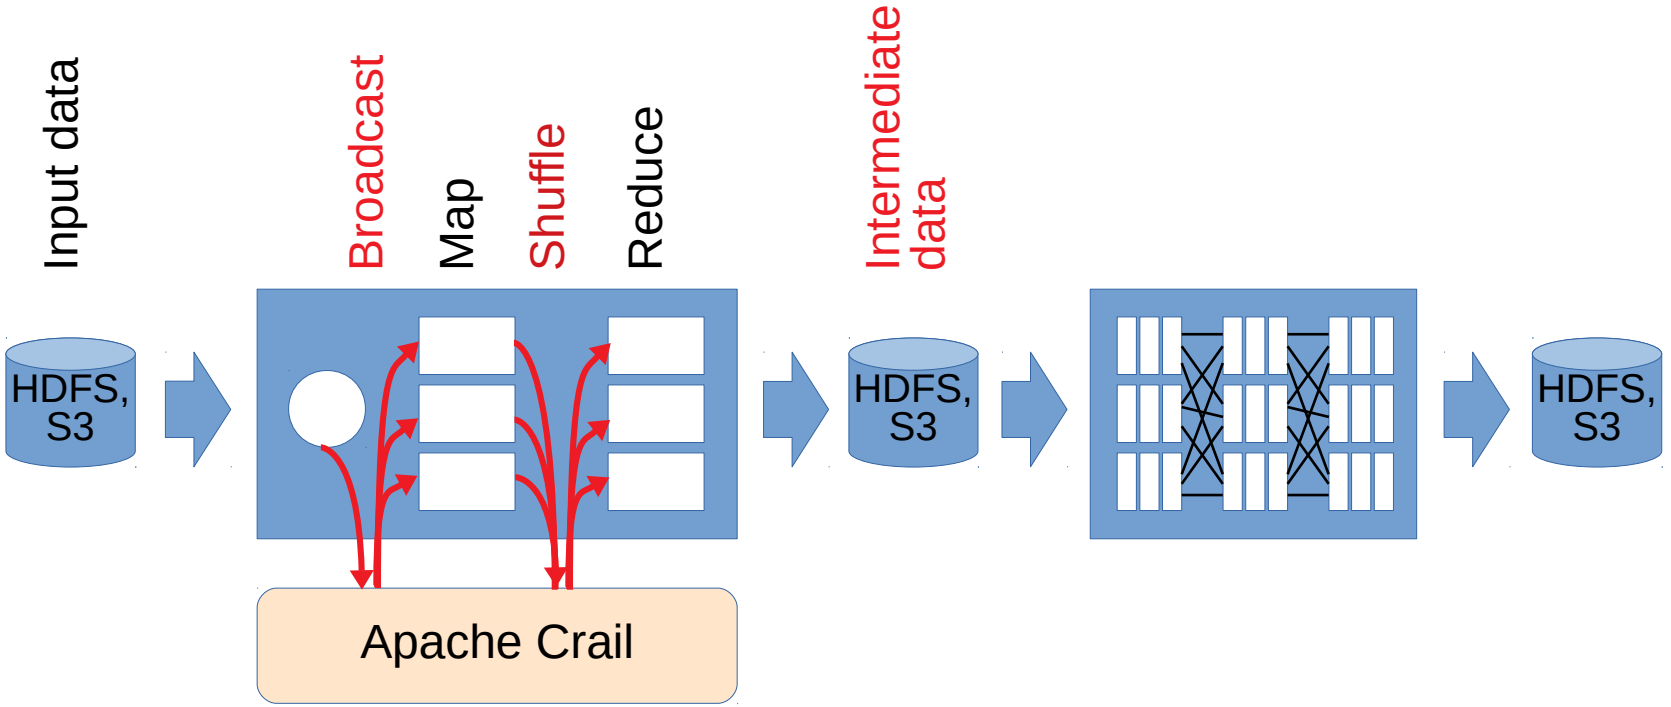

Ephemeral Data

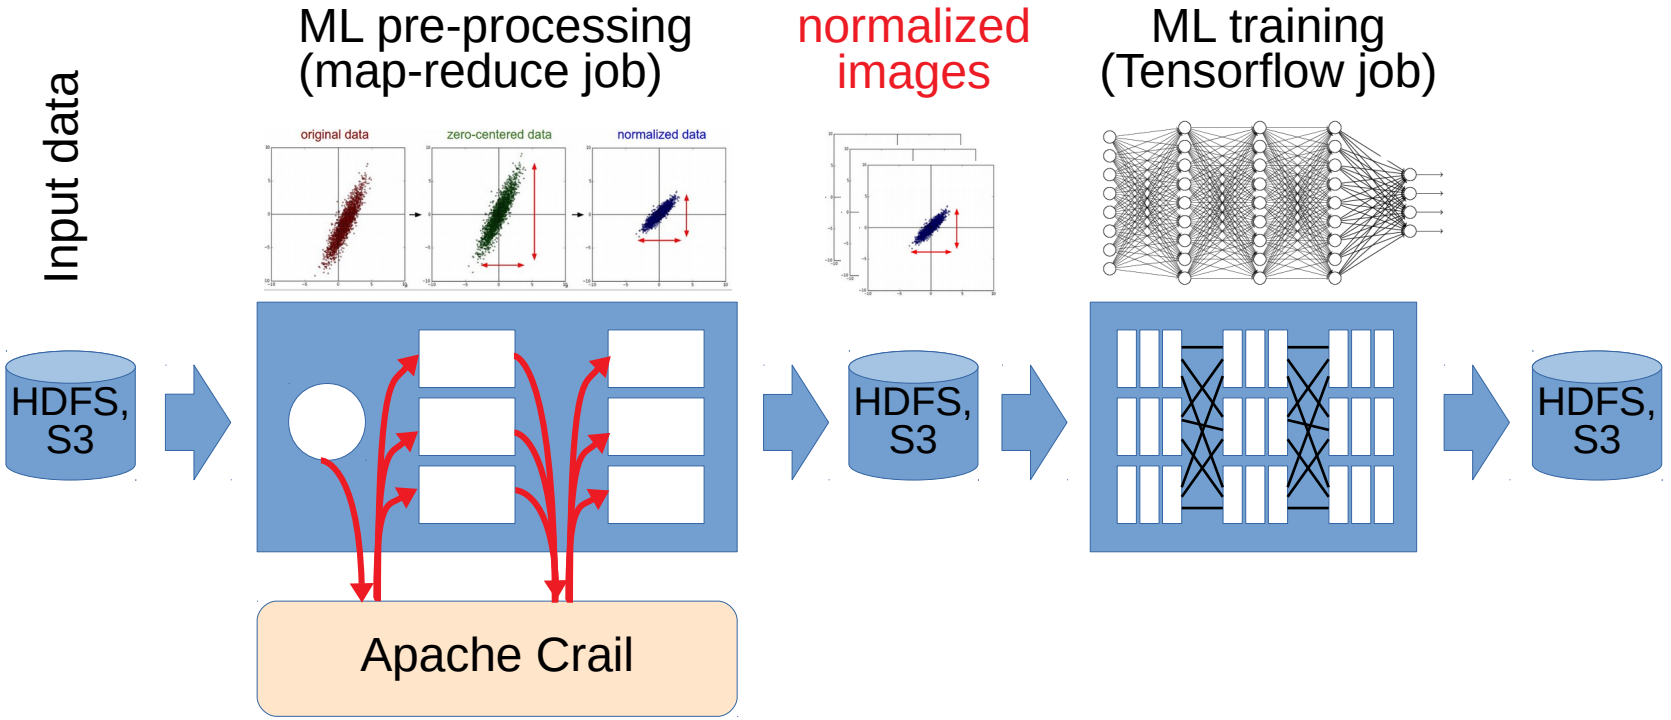

Ephemeral Data

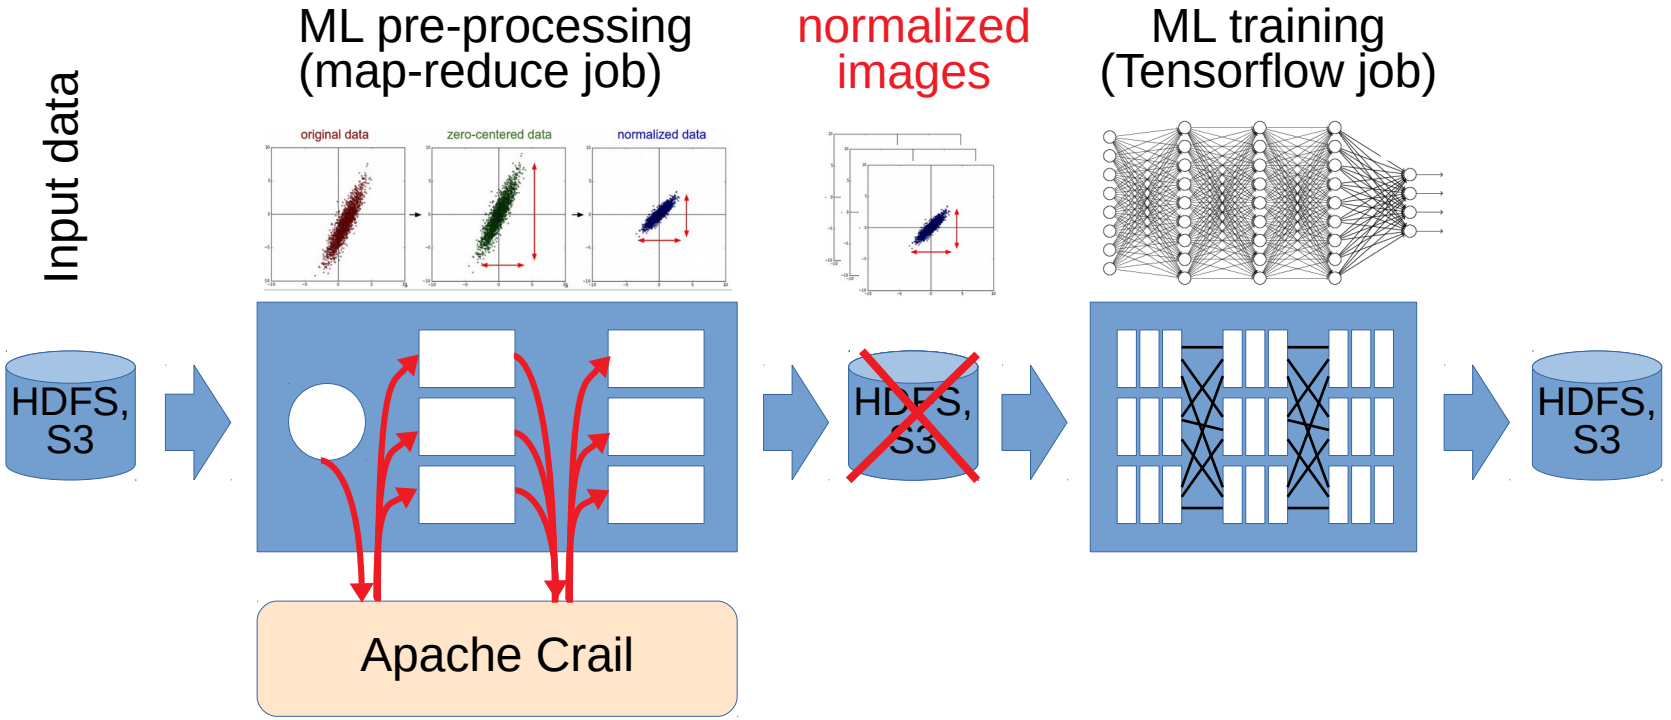

Ephemeral Data

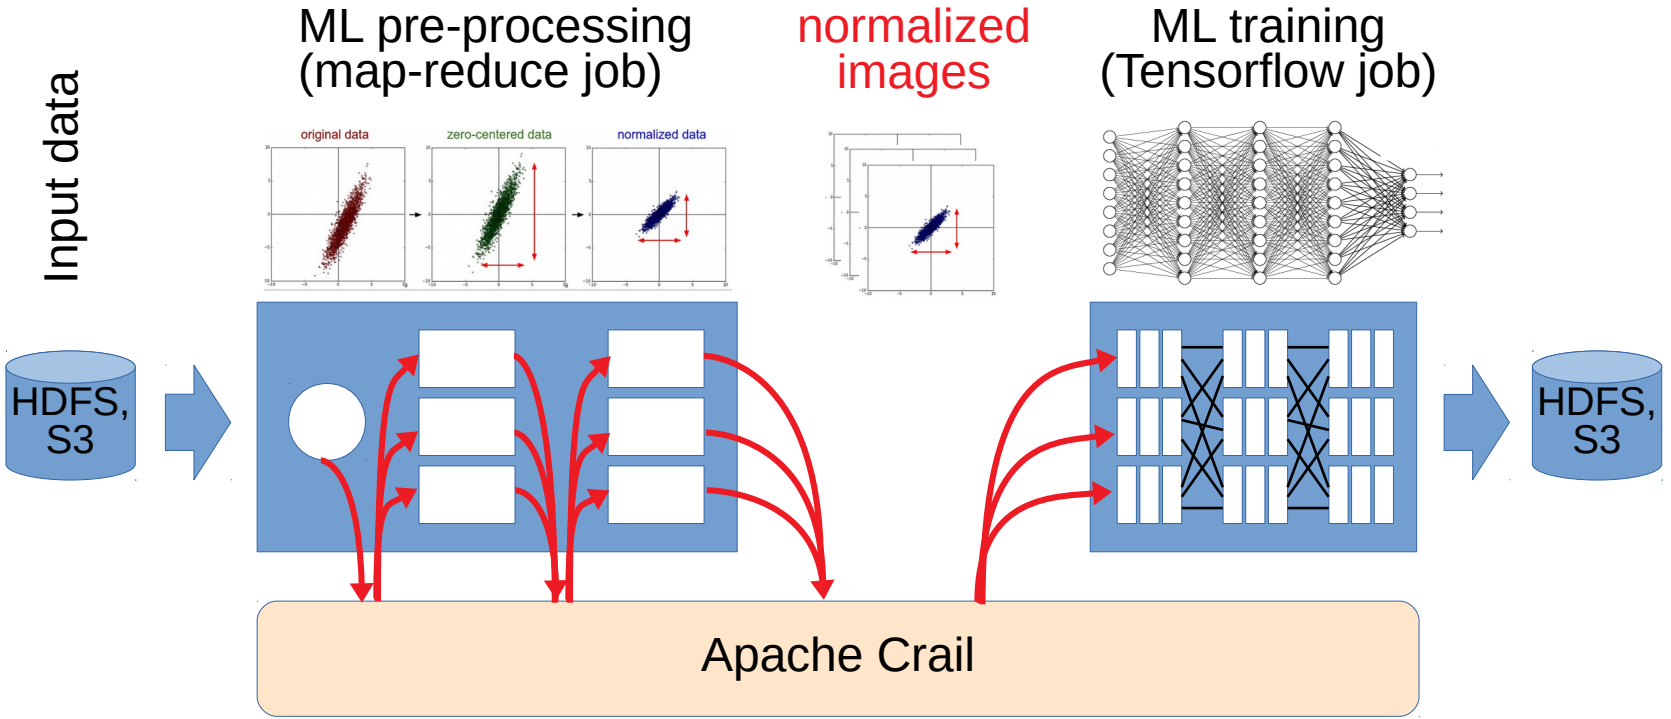

Why/when to use Crail

HDD

10 Gb/s Ethernet

Why/when to use Crail

No Crail needed

HDD

10 Gb/s Ethernet

100MB/s
10ms

10Gb/s
20us

Why/when to use Crail

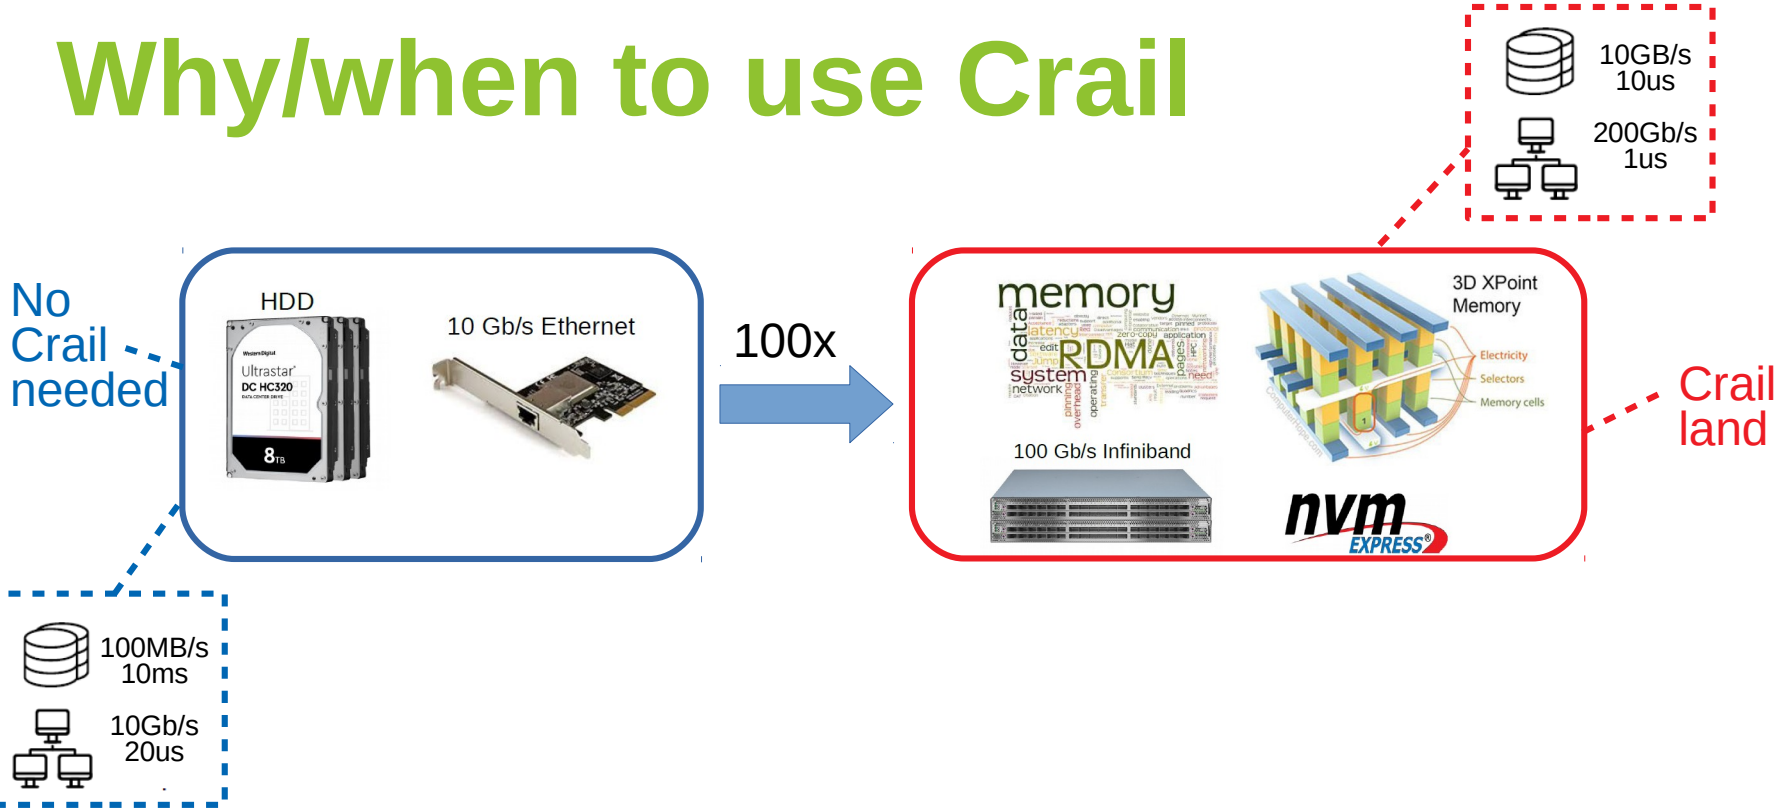

Why/when to use Crail

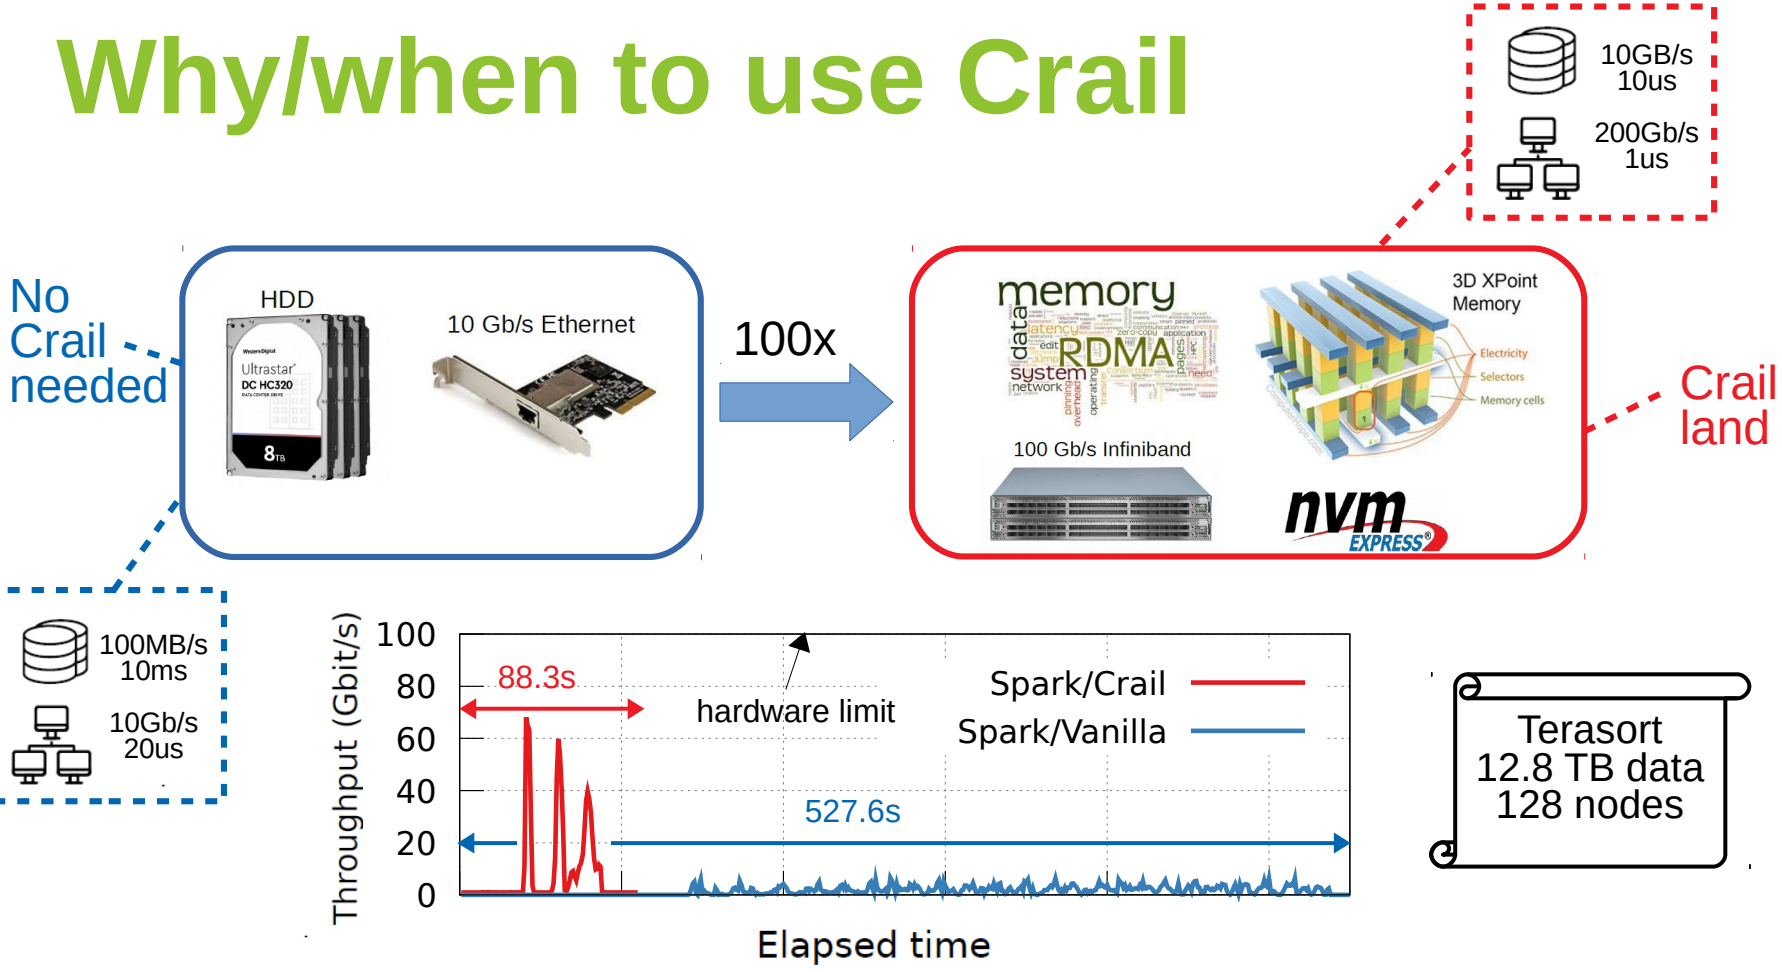

Performance Challenge

	1 Gbps	HDD	100 Gbps	Flash
Bandwidth	117 MB/s	140 MB/s	12.5 GB/s	3.1 GB/s
cycle/unit	38,400	10,957	360	495

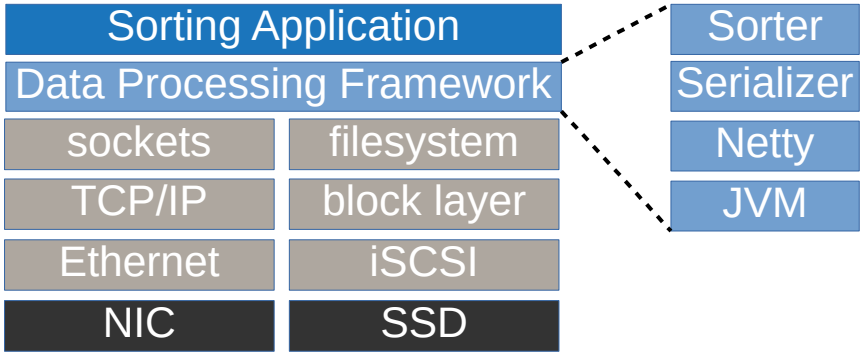

Performance Challenge

	1 Gbps	HDD	100 Gbps	Flash
Bandwidth	117 MB/s	140 MB/s	12.5 GB/s	3.1 GB/s
cycle/unit	38,400	10,957	360	495

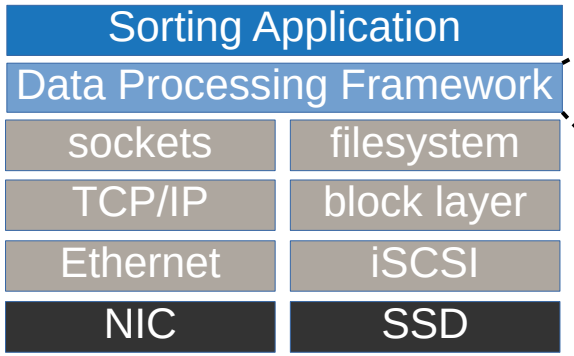

Process chunk
In reduce task

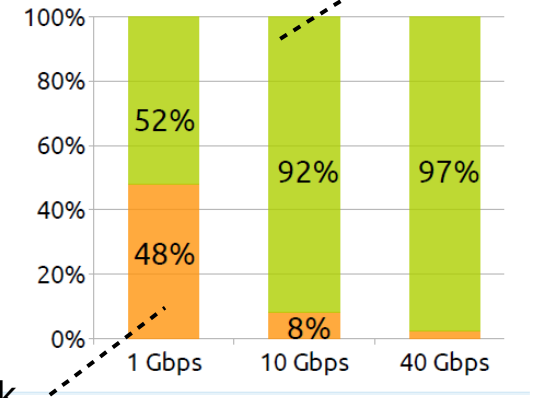

Fetch chunk
Over the network

Performance Challenge

	1 Gbps	HDD	100 Gbps	Flash
Bandwidth	117 MB/s	140 MB/s	12.5 GB/s	3.1 GB/s
cycle/unit	38,400	10,957	360	495

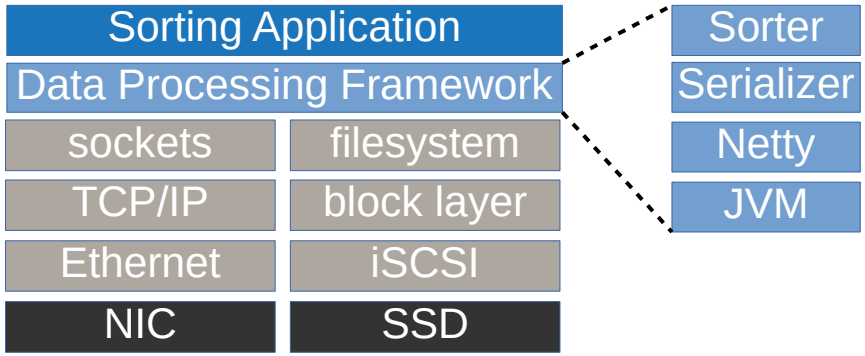

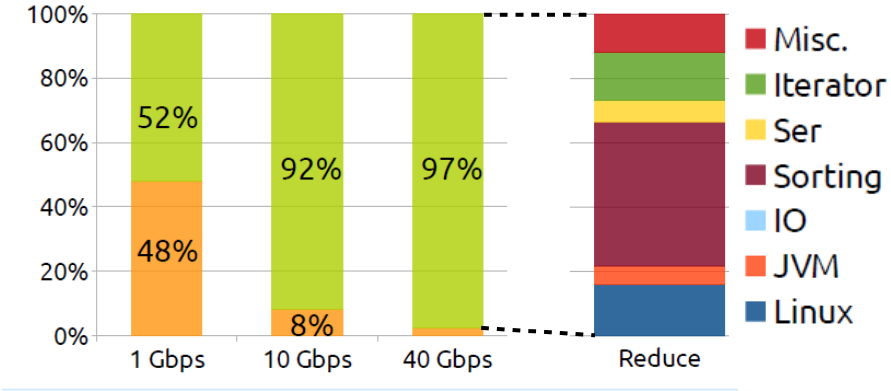

HotNets'16

Performance Challenge

	1 Gbps	HDD	100 Gbps	Flash
Bandwidth	117 MB/s	140 MB/s	12.5 GB/s	3.1 GB/s
cycle/unit	38,400	10,957	360	495

software overhead are spread over the entire stack

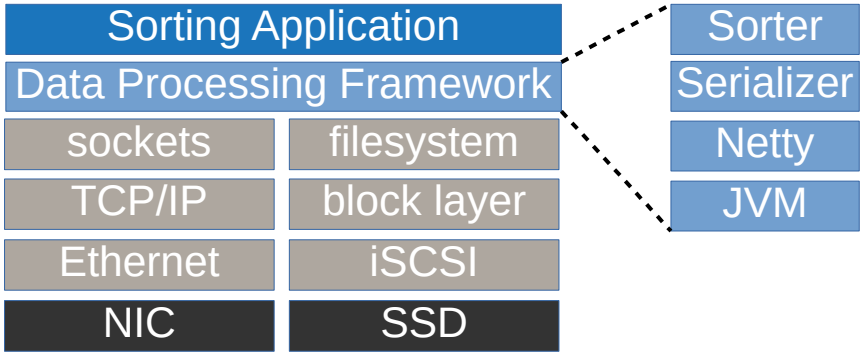

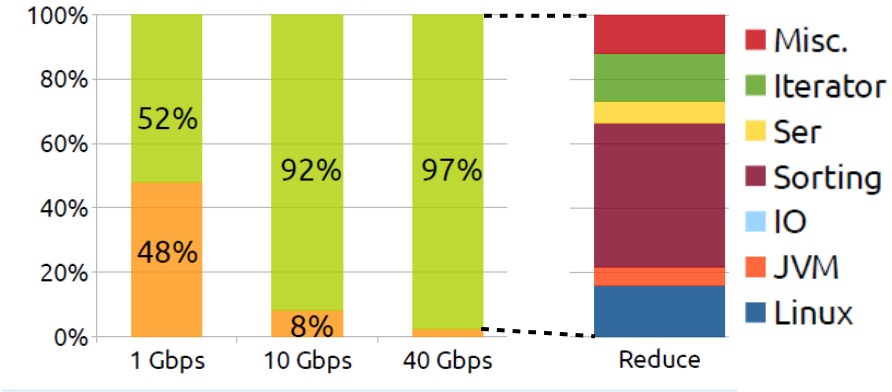

Crail Overview

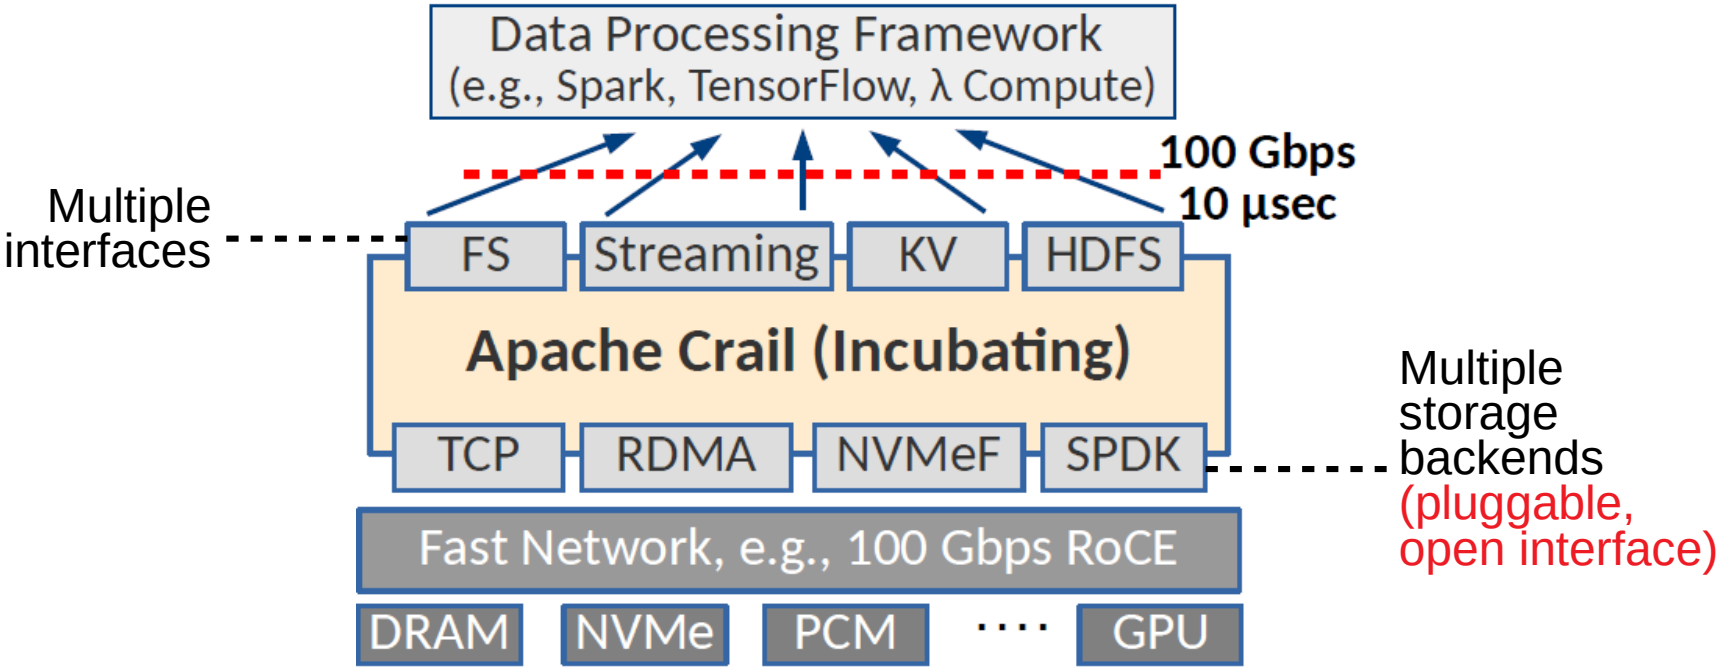

Crail Overview

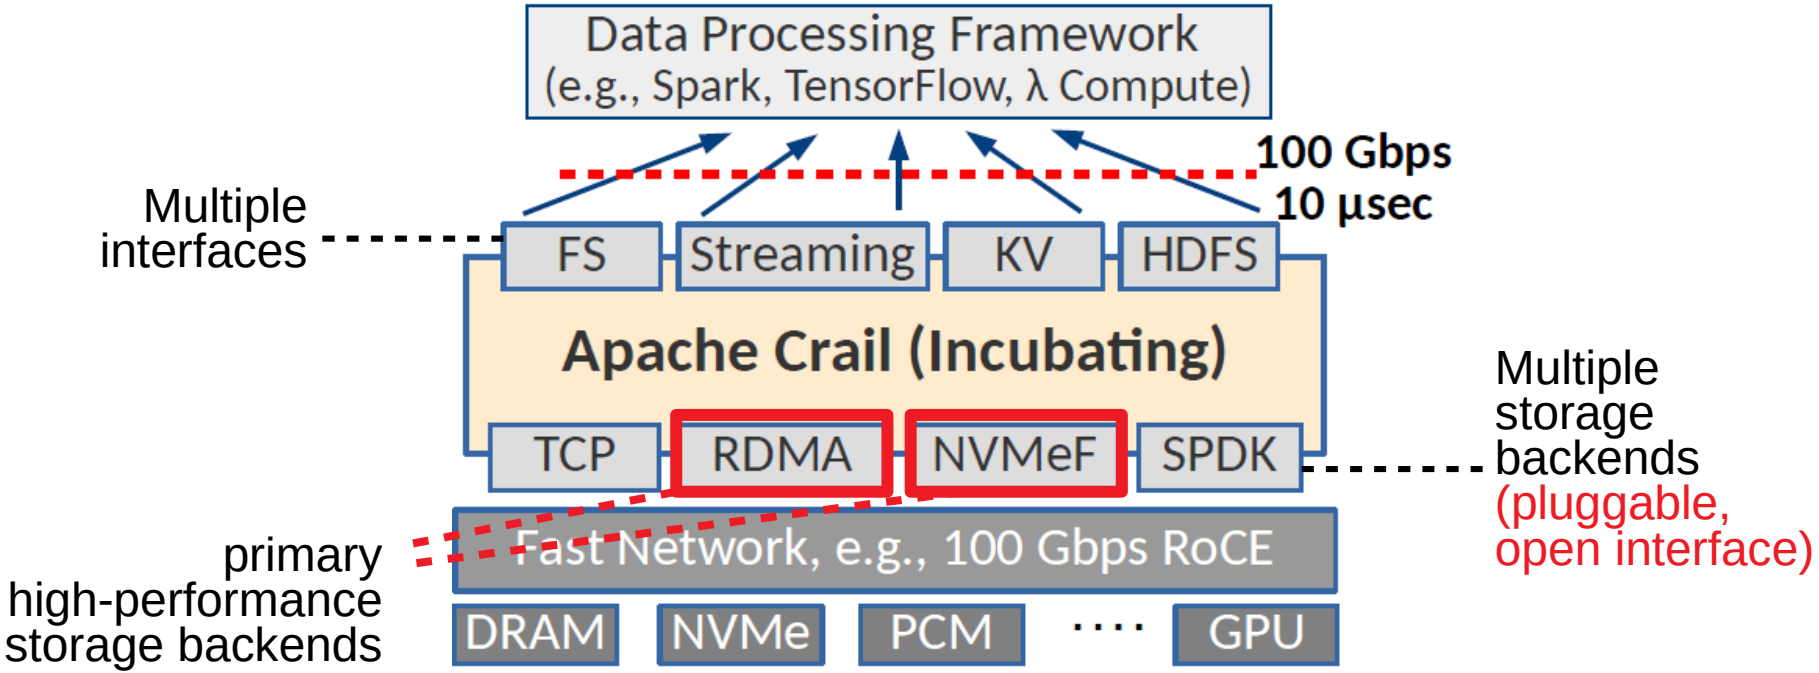

Crail Architecture & API

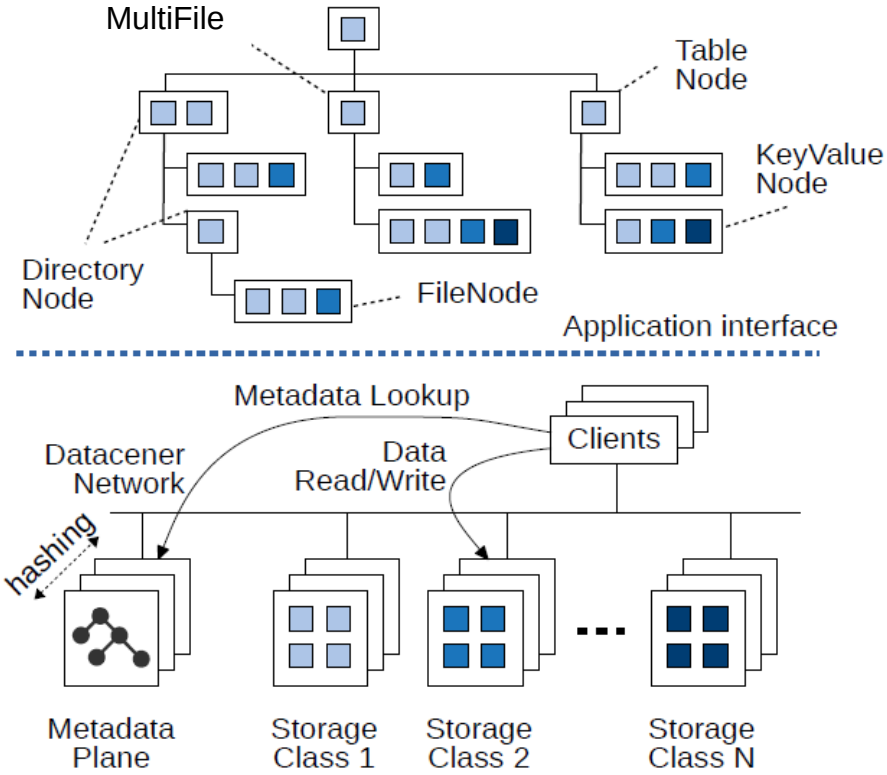

Crail Architecture & API

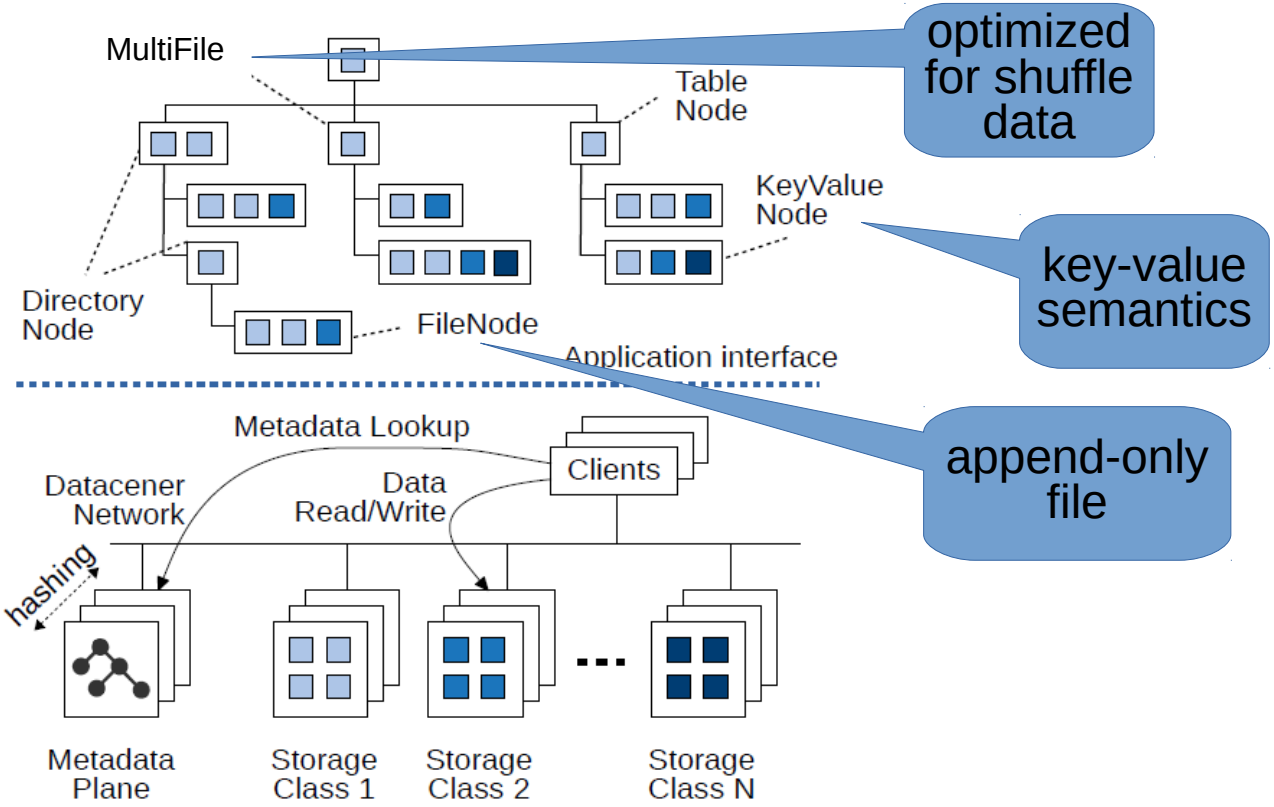

Crail Architecture & API

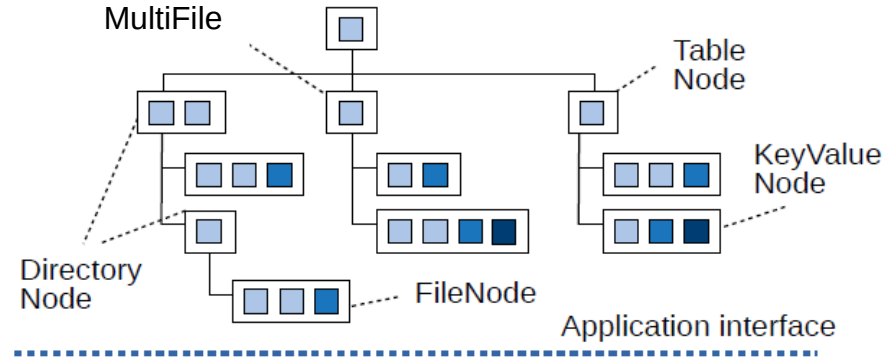

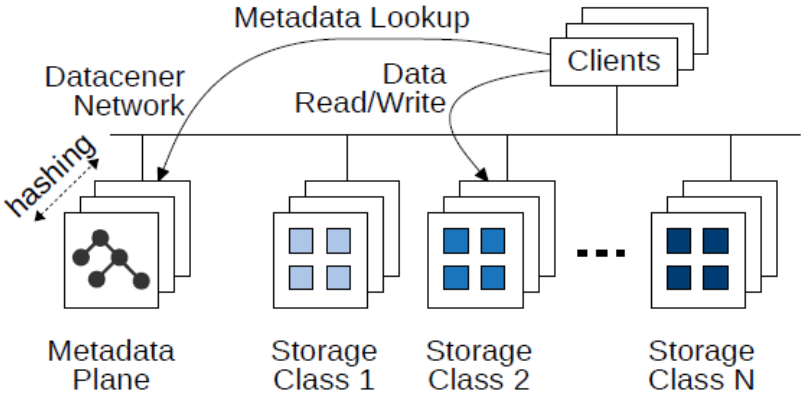

Java:

```
CrailStore crail = CrailStore.newInstance();
Future<Node> fut = crail.create("/a.dat", CrailType.File);
//...do work
CrailFile file = fut.get().asFile();
CrailOutputStream stream = file.getDirectOutputStream();
ByteBuffer buffer = crail.allocateBuffer();
Future<CrailResult> ret = stream.write(buf);
//...do work
ret.get();
```

C++:

```
CrailStore crail;
auto fut = crail.Create<CrailFile>("/tmp.dat");
//..do work
CrailFile file = fut.get();
CrailOutputStream stream = file.outputstream();
shared_ptr<ByteBuffer> buf = make_shared<ByteBuffer>(len);
Future<int> ret = stream.Write(buf);
//..do work
ret.get();
```

Crail Architecture & API

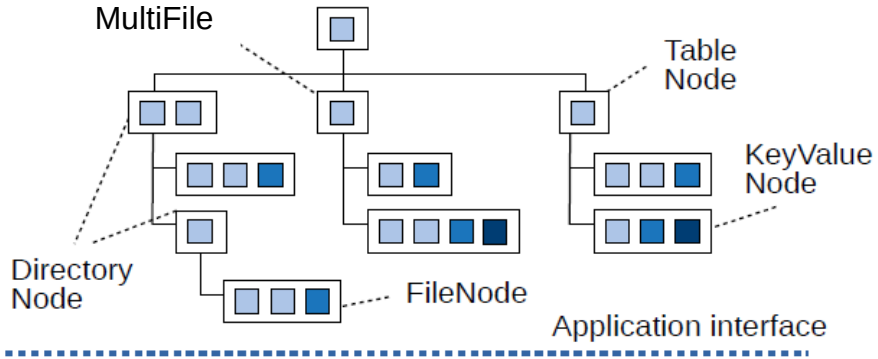

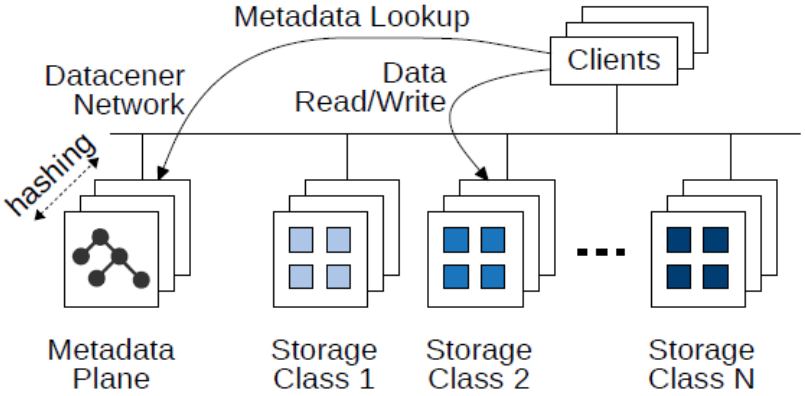

Java:

```
CrailStore crail = CrailStore.newInstance();
Future<Node> fut = crail.create("/a.dat", CrailType.File);
//...do work
CrailFile file = fut.get().asFile();
CrailOutputStream stream = file.getDirectOutputStream();
ByteBuffer buffer = crail.allocateBuffer();
Future<CrailResult> ret = stream.write(buf);
//...do work
ret.get();
```

Node type

C++:

```
CrailStore crail;
auto fut = crail.Create<CrailFile>("/tmp.dat");
//..do work
CrailFile file = fut.get();
CrailOutputStream stream = file.outputstream();
shared_ptr<ByteBuffer> buf = make_shared<ByteBuffer>(len);
Future<int> ret = stream.Write(buf);
//..do work
ret.get();
```

Crail Architecture & API

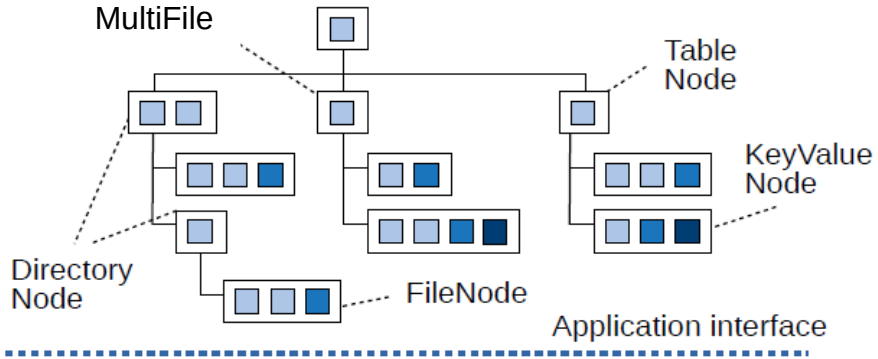

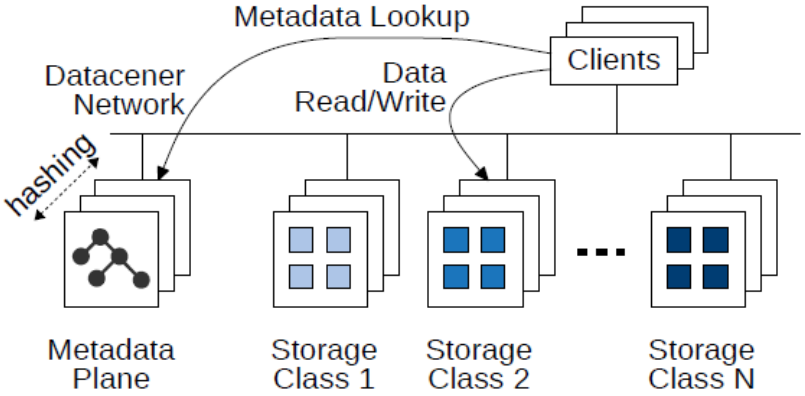

Java:

```
CrailStore crail = CrailStore.newInstance();
Future<Node> fut = crail.create("/a.dat", CrailType.File);
//...do work
CrailFile file = fut.get().asFile();
CrailOutputStream stream = file.getDirectOutputStream();
ByteBuffer buffer = crail.allocateBuffer();
Future<CrailResult> ret = stream.write(buf);
//...do work
ret.get();
```

non-blocking & asynchronous

C++:

```
CrailStore crail;
auto fut = crail.Create<CrailFile>("/tmp.dat");
//..do work
CrailFile file = fut.get();
CrailOutputStream stream = file.outputstream();
shared_ptr<ByteBuffer> buf = make_shared<ByteBuffer>(len);
Future<int> ret = stream.Write(buf);
//..do work
ret.get();
```

Where does the performance come from?

User-Level I/O: Metadata

```
1 auto fut = crail.Create<CrailFile>("/tmp.dat");
  //..do work
2 CrailFile file = fut.get();
```

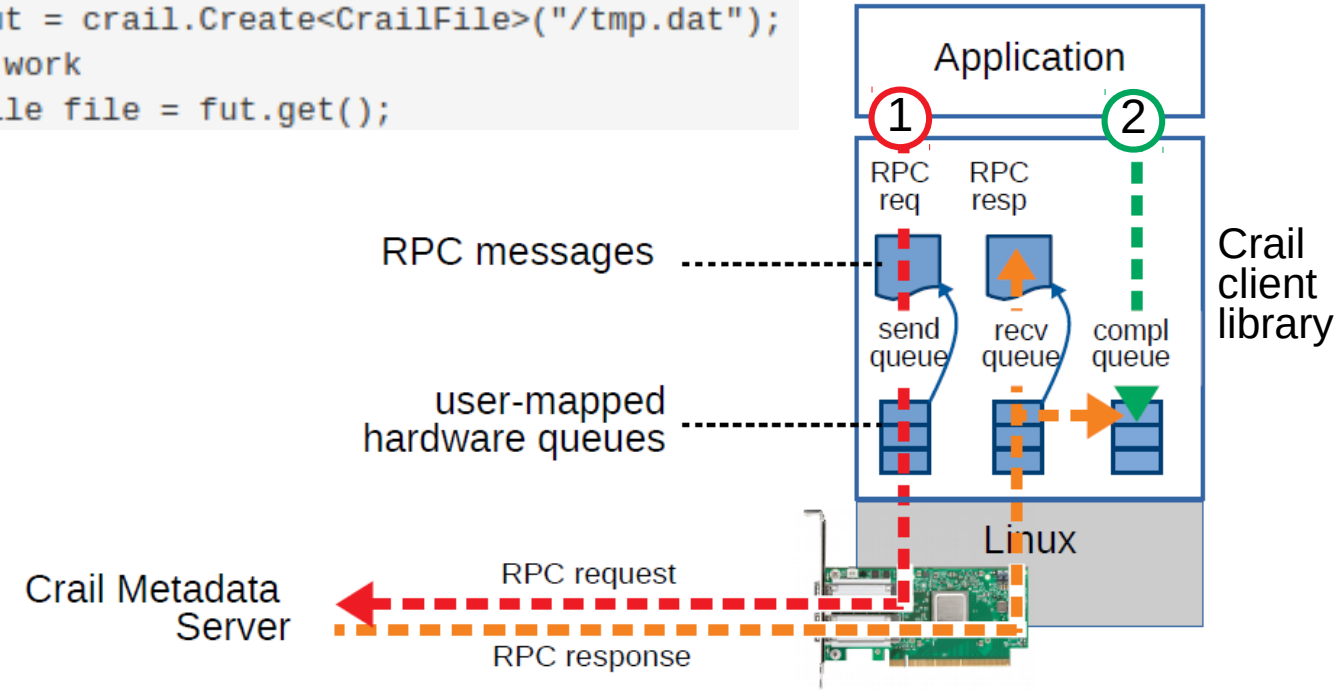

User-Level I/O: Metadata

```
1 auto fut = crail.Create<CrailFile>("/tmp.dat");
  //..do work
2 CrailFile file = fut.get();
```

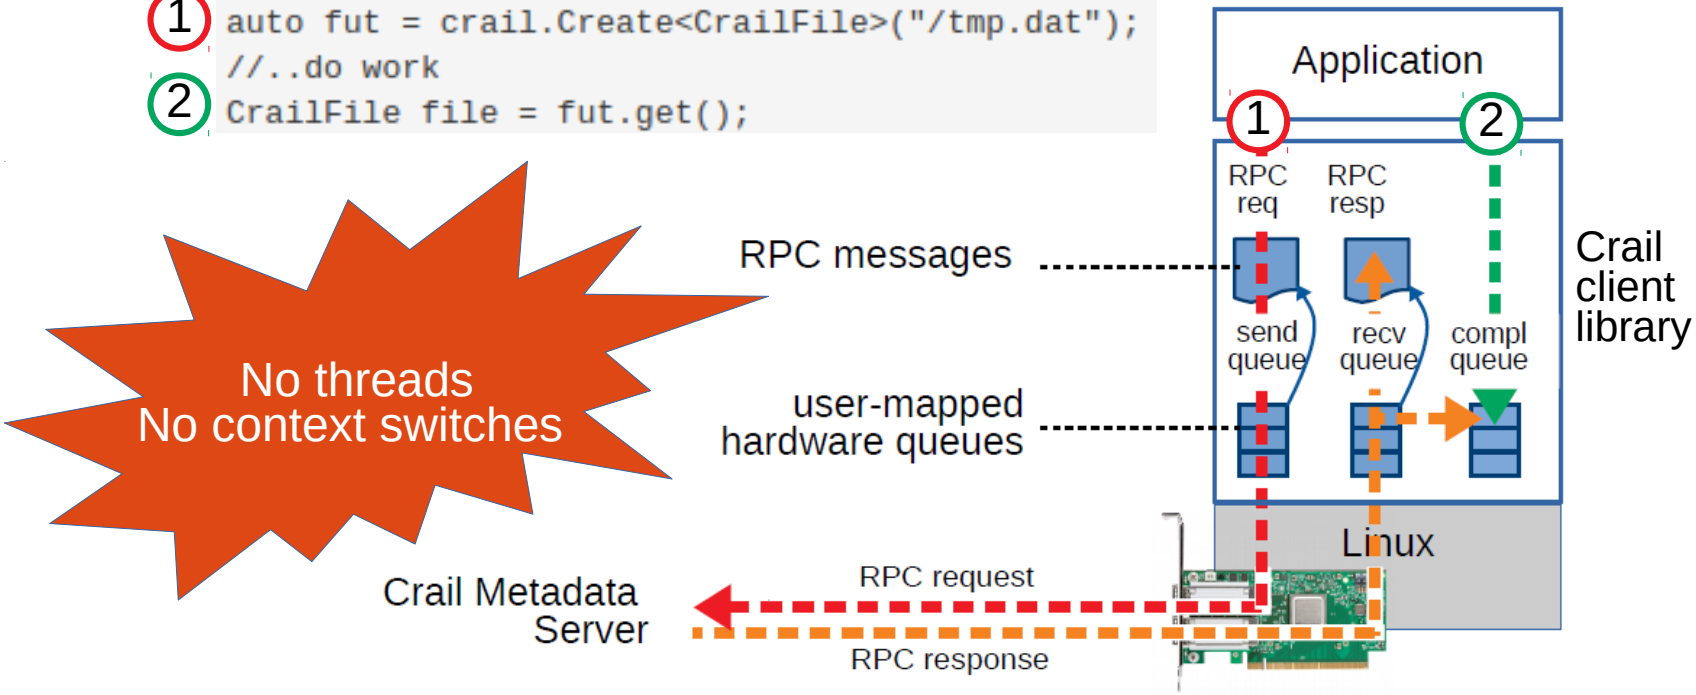

User-Level I/O: Data

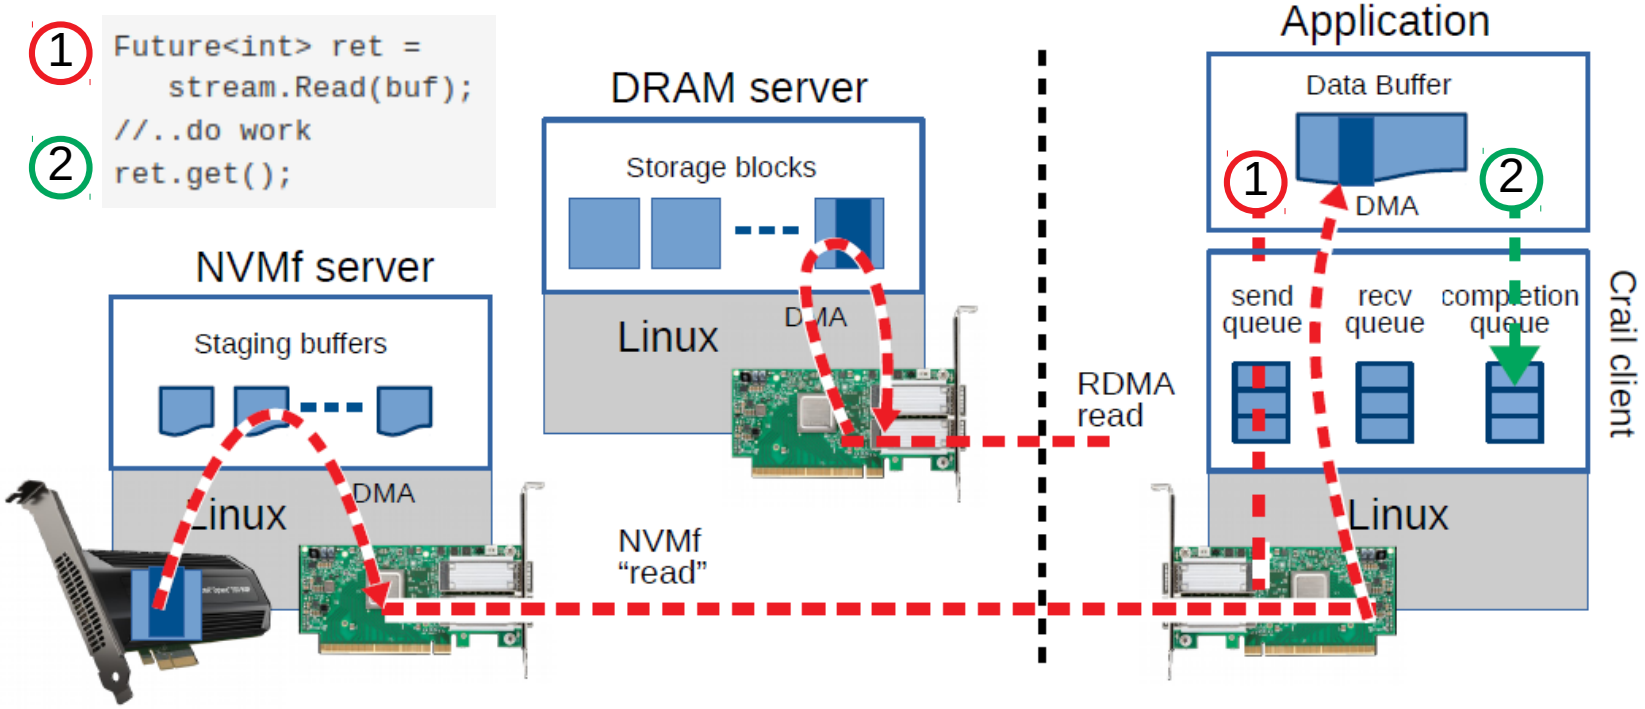

User-Level I/O: Data

zero-copy, transfer only data that is requested

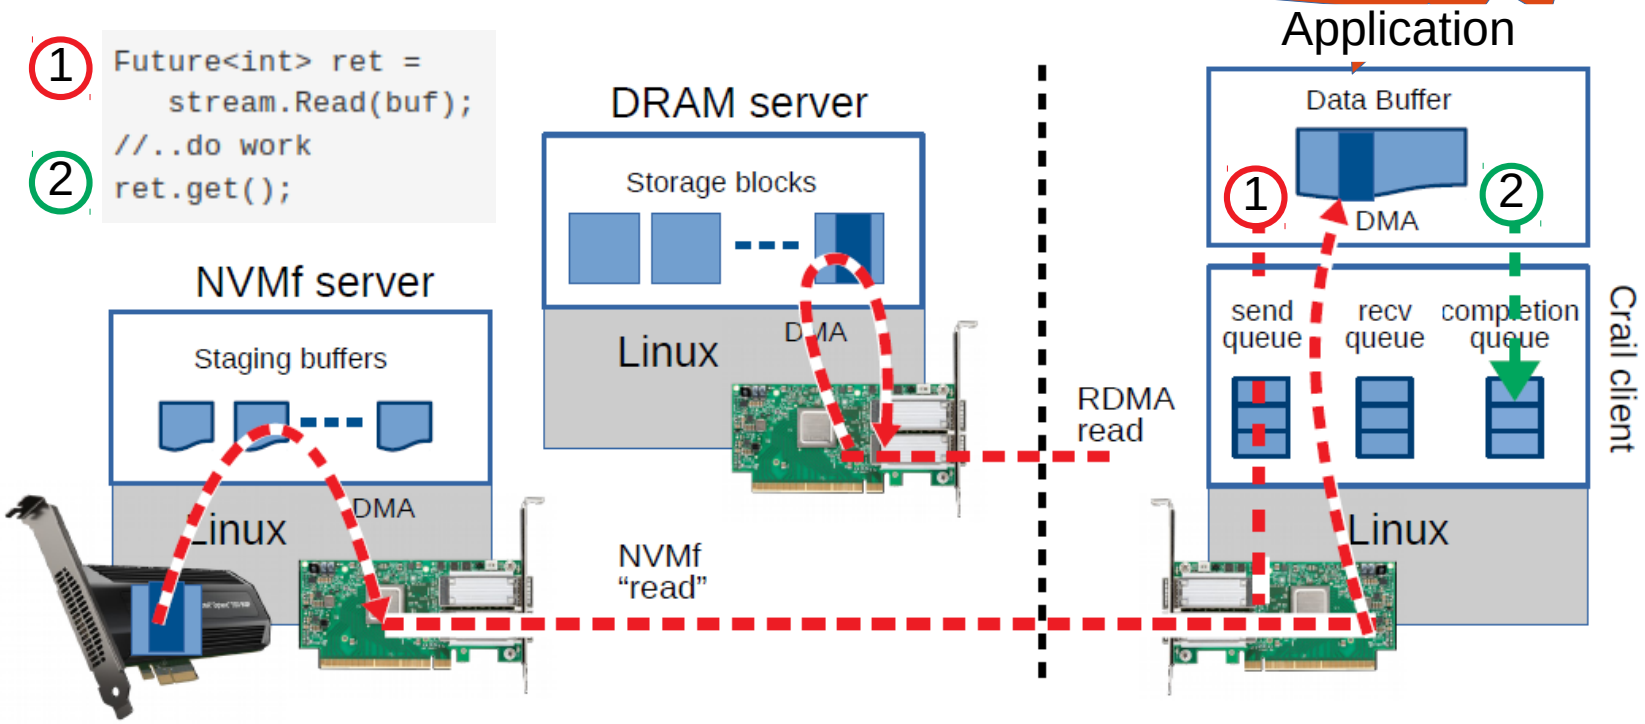

Crail Deployment Modes

compute/storage
co-located

storage
disaggregation

flash storage
disaggregation

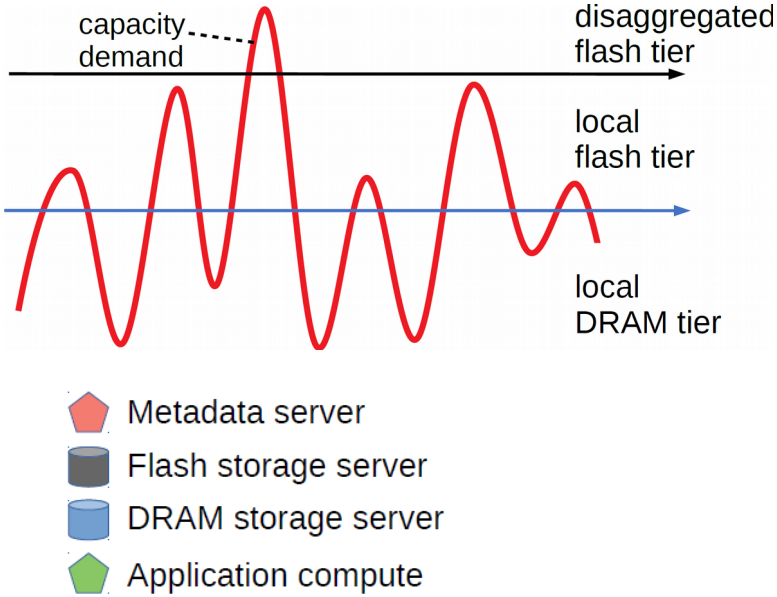

YCSB KeyValue Workload

GET
Value size:
1KB

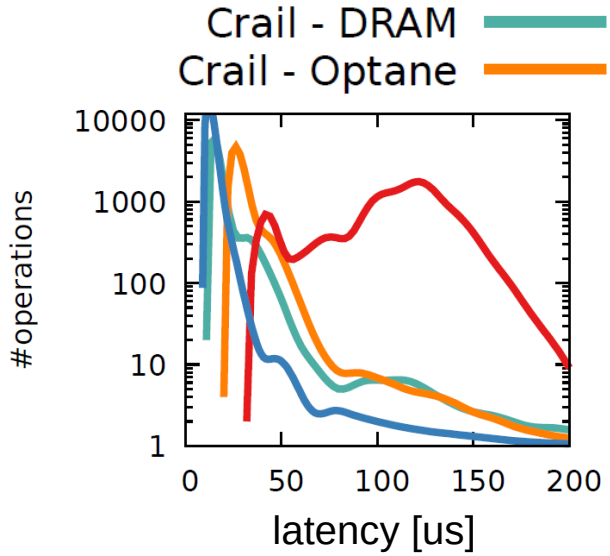

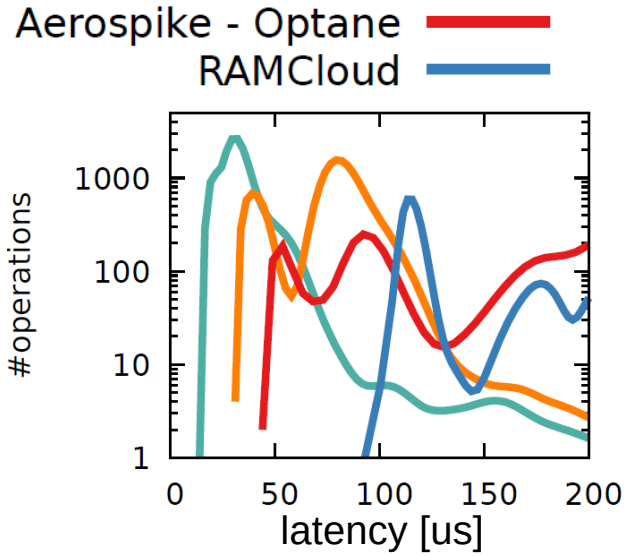

GET
Value size:
100KB

Crail offers Get latencies of ~12us and 30us for DRAM and NVM for 100 byte KV pairs
Crail offers Get latencies of ~30us and 40us for DRAM and NVM for 1000 byte KV pairs

Spark GroupBy (80M keys, 4K)

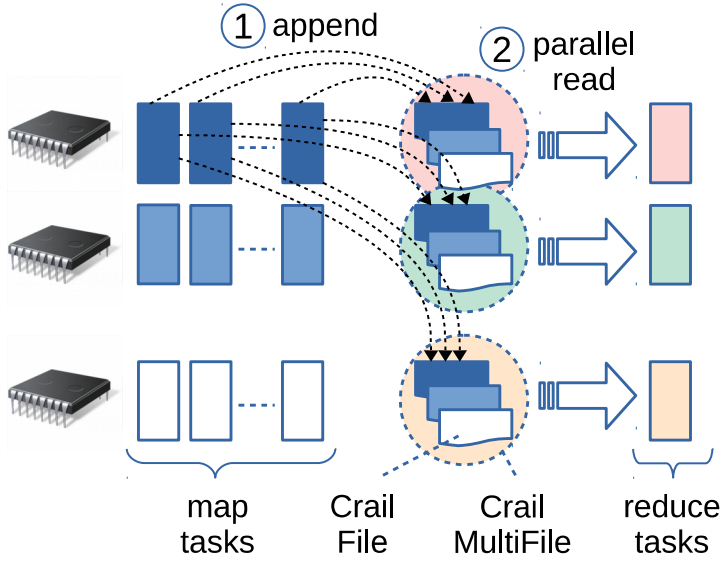

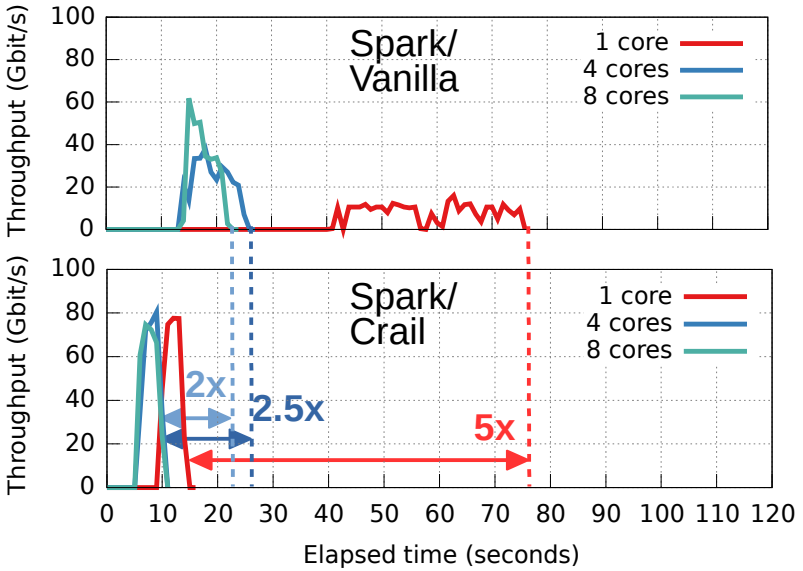

Spark shuffling via Crail on a single core is 2x faster
than vanilla Spark on 8 cores per executor (8 executors)

DRAM & Flash Disaggregation

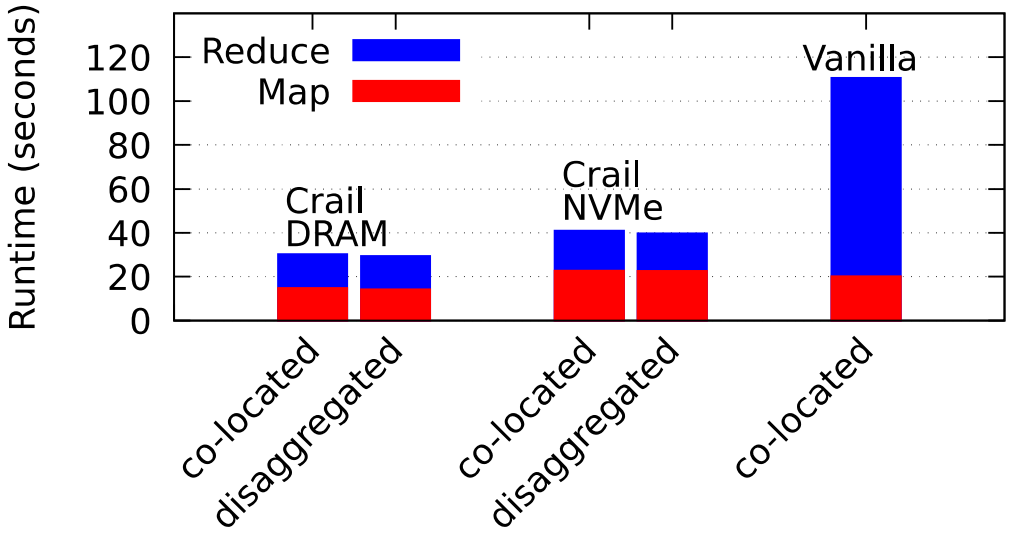

Crail enables disaggregation of temporary data at no cost

DRAM/Flash Tiering

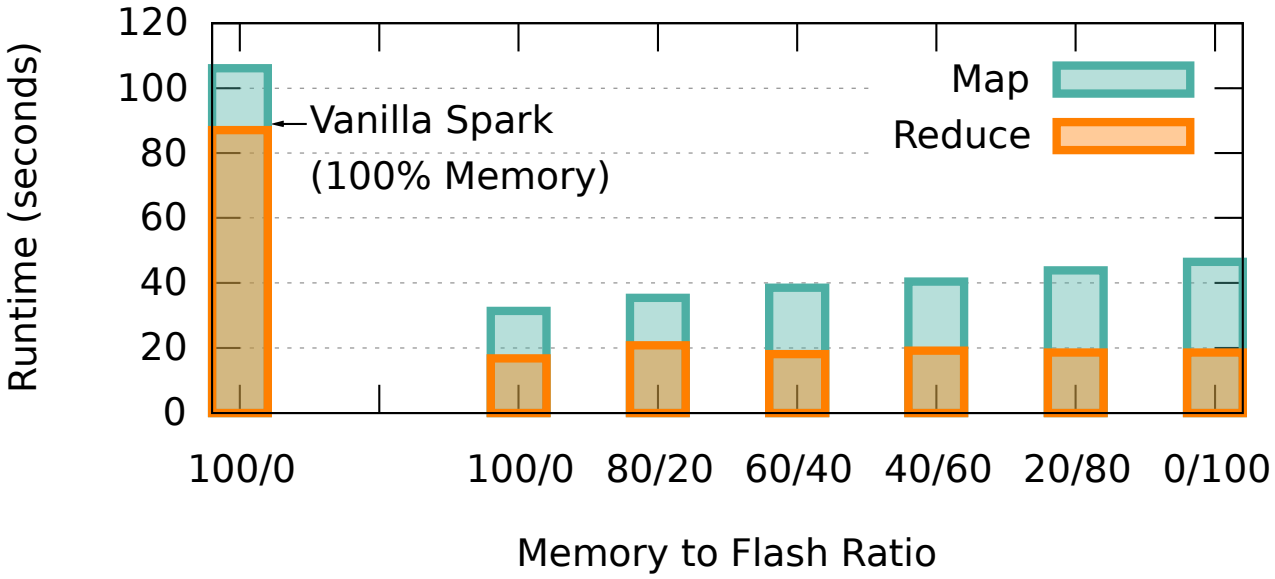

Using flash only increases the sorting time by around 48%

Conclusions

- Apache Crail: Fast distributed “tmp”
  - User-level I/O
  - Storage disaggregation
  - Memory/flash convergence
- Applications
  - Intra-job scratch space (shuffle, broadcast, etc.)
  - Multi-job pipelines
- Coming soon
  - Native Crail (C++)
  - Tensorflow-Crail