

Strata Data 2019
San Francisco, CA



Tuning Performance for SQL-on-Anything Analytics



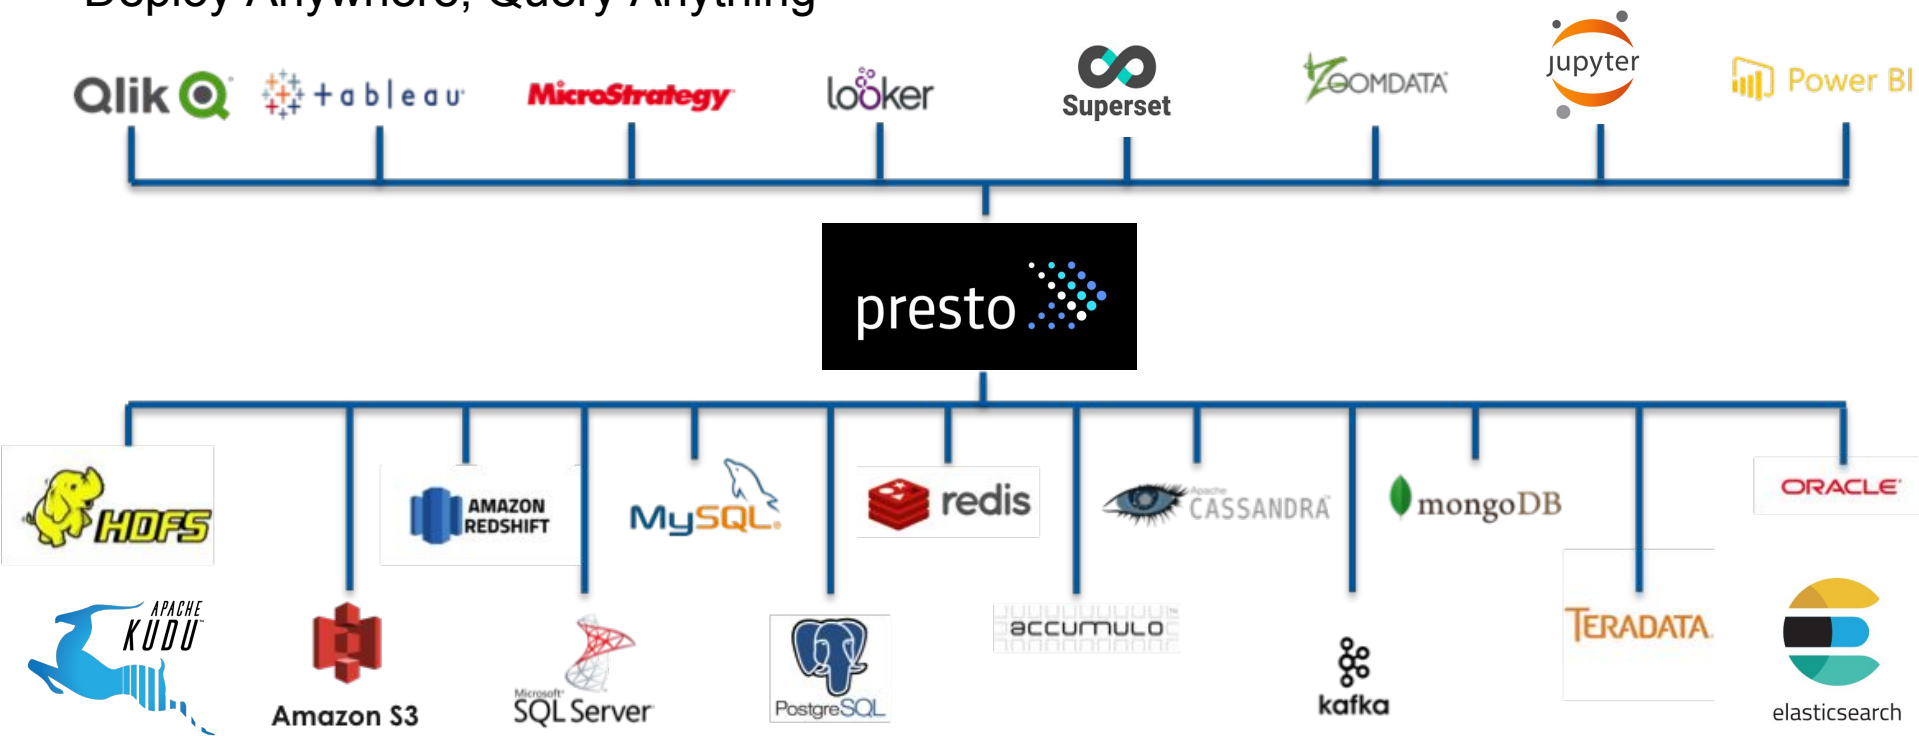
@prestosql @starburstdata

Martin Traverso, Co-creator of Presto

Kamil Bajda-Pawlikowski, CTO Starburst

Presto: SQL-on-Anything

Deploy Anywhere, Query Anything



Why Presto?



Community-driven
open source project



High performance ANSI SQL engine

- New Cost-Based Query Optimizer
- Proven scalability
- High concurrency



Separation of compute
and storage

- Scale storage and compute independently
- No ETL or data integration necessary to get to insights
- SQL-on-anything



No vendor lock-in

- No Hadoop distro vendor lock-in
- No storage engine vendor lock-in
- No cloud vendor lock-in

Project History



FALL 2012

4 developers
start Presto
development

FALL 2013

Facebook open
sources Presto

SPRING 2015

Teradata joins the
community, begins
investing heavily in
the project

SUMMER 2017

180+ Releases
50+ Contributors
5000+ Commits

WINTER 2017

Starburst is founded by
a team of Presto
committers, Teradata
veterans

WINTER 2019

Presto Software
Foundation
established

Community

See more at [our Wiki](#)



Presto in Production

Facebook: 10,000+ of nodes, HDFS (ORC, RCFile), sharded MySQL, 1000s of users

Uber: 2,000+ nodes (several clusters on premises) with 160K+ queries daily over HDFS (Parquet/ORC)

Twitter: 2,000+ nodes (several clusters on premises and GCP), 20K+ queries daily (Parquet)

LinkedIn: 500+ nodes, 200K+ queries daily over HDFS (ORC), and ~1000 users


Lyft: ----- *redacted due to the quiet period for the IPO* -----

Netflix: 300+ nodes in AWS, 100+ PB in S3 (Parquet)

Yahoo! Japan: 200+ nodes for HDFS (ORC), and ObjectStore

FINRA: 120+ nodes in AWS, 4PB in S3 (ORC), 200+ users

Starburst Data



Founded by Presto committers:

- Over 4 years of contributions to Presto
- Presto distro for on-prem and cloud env
- Supporting large customers in production
- Enterprise subscription add-ons (ODBC, Ranger, Sentry, Oracle, Teradata)

Notable features contributed:

- ANSI SQL syntax enhancements
- Execution engine improvements
- Security integrations
- Spill to disk
- Cost-Based Optimizer

<https://www.starburstdata.com/presto-enterprise/>

Performance

Built for Performance

Query Execution Engine:

- MPP-style **pipelined** in-memory execution
- **Columnar** and **vectorized** data processing
- Runtime query **bytecode compilation**
- Memory **efficient** data structures
- Multi-threaded multi-core execution
- Optimized readers for **columnar formats** (ORC and Parquet)
- Predicate and column projection **pushdown**
- Now also **Cost-Based Optimizer**

CBO in a nutshell

Presto Cost-Based Optimizer includes:

- support for **statistics** stored in Hive Metastore
- **join reordering** based on selectivity estimates and cost
- automatic **join type** selection (repartitioned vs broadcast)
- automatic left/right side selection for joined tables

<https://www.starburstdata.com/technical-blog/>

Statistics & Cost


Hive Metastore statistics:

- number of rows in a table
- number of distinct values in a column
- fraction of NULL values in a column
- minimum/maximum value in a column
- average data size for a column

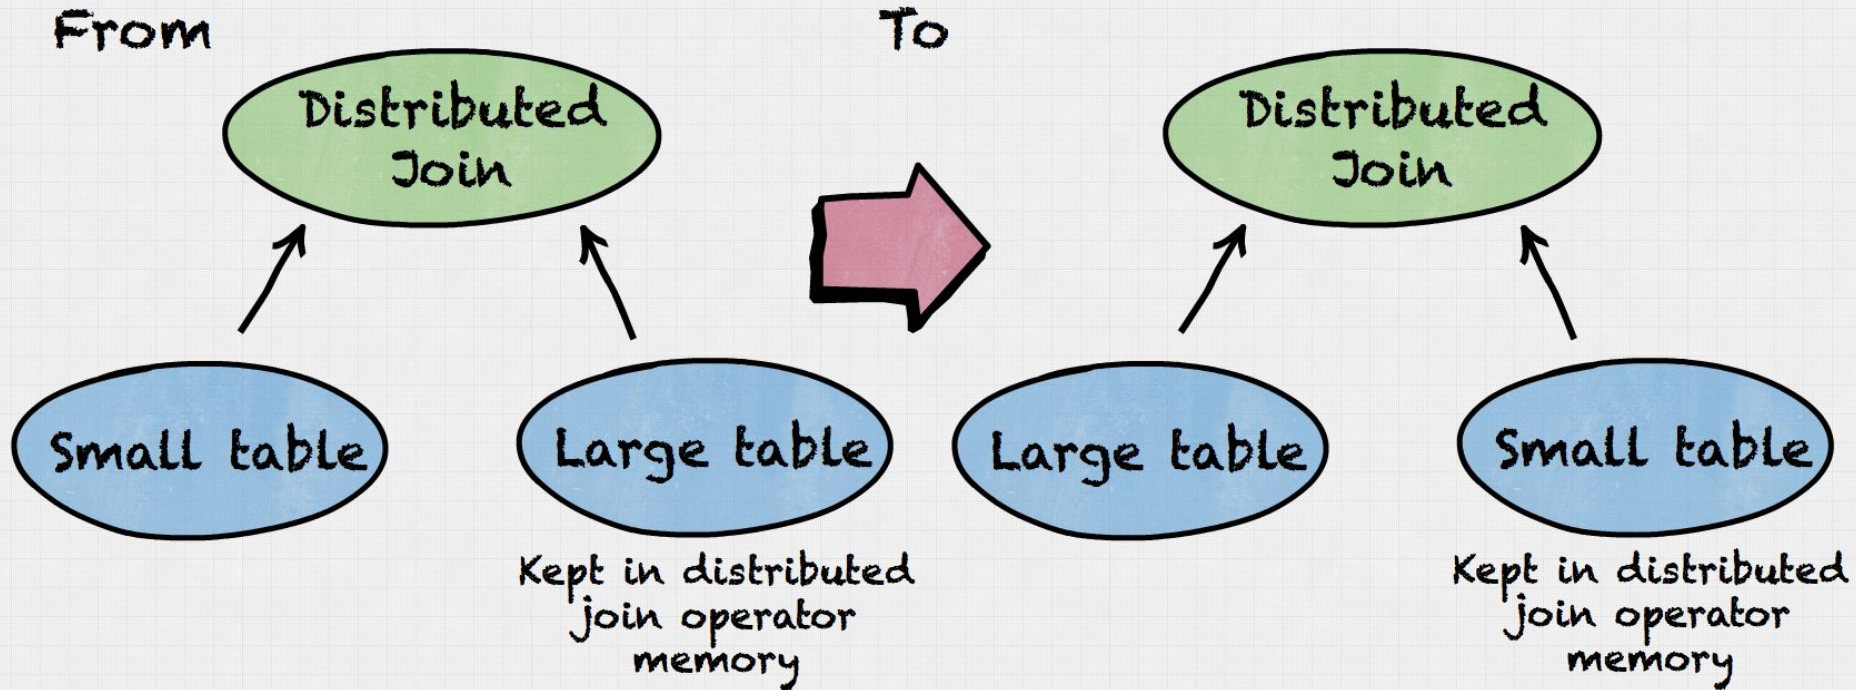
Cost calculation includes:

- CPU
- Memory
- Network I/O

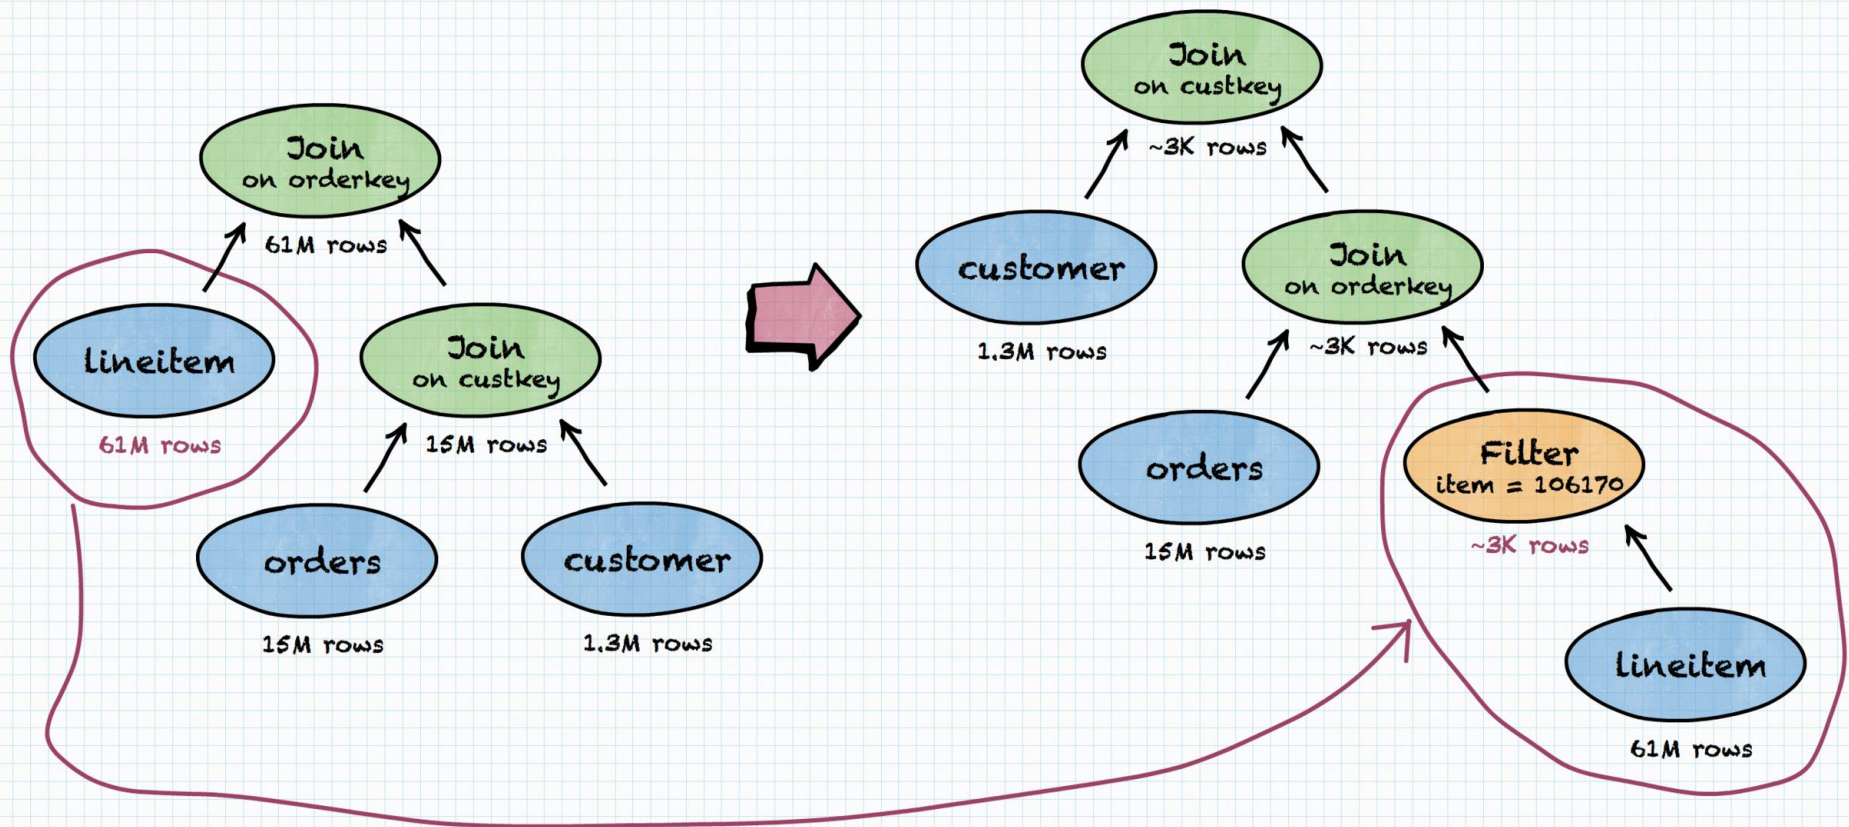
Join type selection



Join left/right side decision




Join reordering with filter

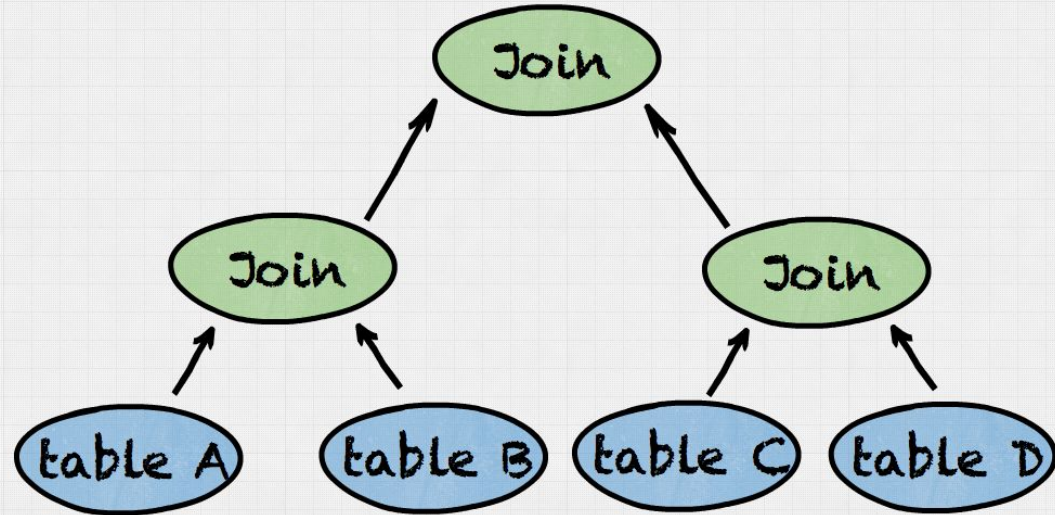


Join tree shapes


Left deep



Bushy tree



Benchmark results




<https://www.starburstdata.com/presto-benchmarks/>


Benchmark results

- on average 7x improvement vs EMR Presto
- EMR Presto cannot execute many TPC-DS queries
- All TPC-DS queries pass on Starburst Presto

Starburst Presto (CBO) vs EMR Presto speedup



EMR Presto TPC-DS passed queries %



<https://www.starburstdata.com/presto-aws/>

Recent CBO enhancements

- Deciding on semi-join distribution type based on cost
- Support for outer joins
- Capping a broadcasted table size
- Various minor fixes in cardinality estimation
- ANALYZE table (native in Presto)
- Stats for AWS Glue Catalog (exclusive from Starburst)

Current and Future work

What's next for Optimizer

- Stats support
 - Improved stats for Hive
 - Stats for DBMS connectors and NoSQL connectors
 - Tolerate missing / incomplete stats
- Core CBO enhancements
 - Cost more operators
 - Adjust cost model weights based on the hardware
 - Adaptive optimizations
 - Introduce Traits
- Involve connectors in optimizations

Involving Connectors in Optimization

History and Current State

- Original motivation: partition pruning for queries over Hive tables
- Simple range predicates and nullability checks passed to connectors.

Modeled as [TupleDomain](#)

```
((col0 BETWEEN ? AND ?) OR (col0 BETWEEN ? and ?) OR ...))  
AND  
((col1 BETWEEN ? AND ?) OR (col1 BETWEEN ? and ?) OR ...))  
AND  
...
```

History and Current State

- Partial evaluation of non-trivial expressions
 - Bind only known variables
 - Result in "true/false/null" or "can't tell". E.g.,

$f(a, b) := \text{lower}(a) \text{ LIKE 'john\%' AND } b = 1$

$f(\text{'Mary'}, ?) \rightarrow \text{false} \rightarrow \text{can prune}$




$f(\text{'John S'}, ?) \rightarrow b = 1 \rightarrow \neg_(\text{ツ})_/\neg$

Beyond Simple Filter Pushdown...




- Dereference expressions. E.g., `x.a > 5`
- Array/map subscript. E.g., `a['key'] = 10`
- Complex filters and projections
- Aggregations
- Joins
- Limit: <https://github.com/prestosql/presto/pull/421>
- Sampling
- Others...

<https://github.com/prestosql/presto/issues/18>

Rule 1

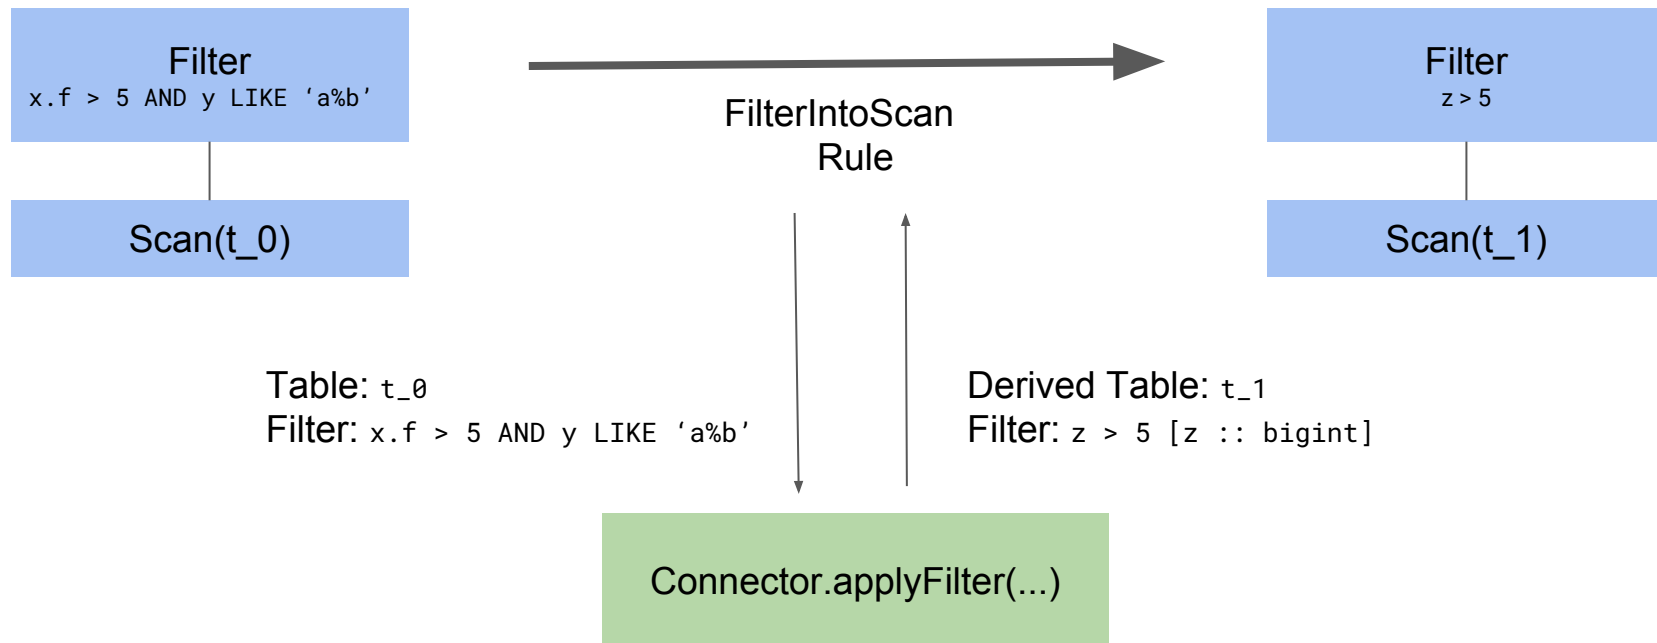


Rule 2



```
SELECT count(*)  
FROM t  
WHERE x.f > 5 AND y LIKE 'a%b'
```

```
Table t  
x :: row(f bigint, g bigint)  
y :: varchar(10)
```



New Connector APIs

applyFilter(ConnectorTableHandle table, Expression filter)

applyLimit(ConnectorTableHandle table, long limit)

applyAggregation(ConnectorTableHandle table, List<Aggregation> aggregates)


applySampling(ConnectorTableHandle table, double samplingRate)

...

Performance Benefits (?)


- Better support for sophisticated backend systems
 - Druid, Pinot, ElasticSearch
 - SQL databases
- Improved performance for columnar data formats (Parquet, ORC)

ORC Performance Improvements



<https://github.com/prestosql/presto/pull/555>

ORC Performance Improvements - TPC-DS



Project Roadmap

- Coordinator HA
- Kubernetes
- Dynamic filtering
- Connectors
 - Phoenix
 - Iceberg
 - Druid
- TIMESTAMP semantics
- And more... <https://github.com/prestosql/presto/labels/roadmap>

Getting Involved

- Join us on Slack
 - Invite link: <https://prestosql.io/community.html>
- Github: <https://github.io/prestosql/presto>
- Website: <https://prestosql.io>

Further reading

<https://www.starburstdata.com/presto-newsletter/>

<https://fivetran.com/blog/warehouse-benchmark>

<https://www.concurrencylabs.com/blog/starburst-presto-vs-aws-emr-sql/>

<http://bytes.schibsted.com/bigdata-sql-query-engine-benchmark/>

<https://virtuslab.com/blog/benchmarking-spark-sql-presto-hive-bi-processing-googles-cloud-dataproc/>

Thank You!



www.prestosql.io



Starburst

www.starburstdata.com



@prestosql @starburstdata